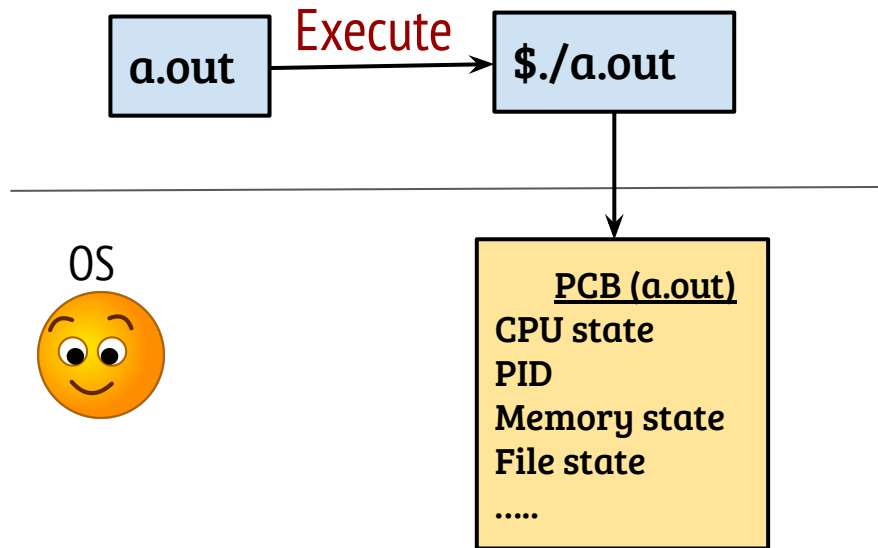


CS330: Operating Systems

Process API: System calls

Recap: The process abstraction


- The OS creates a process when we run an executable



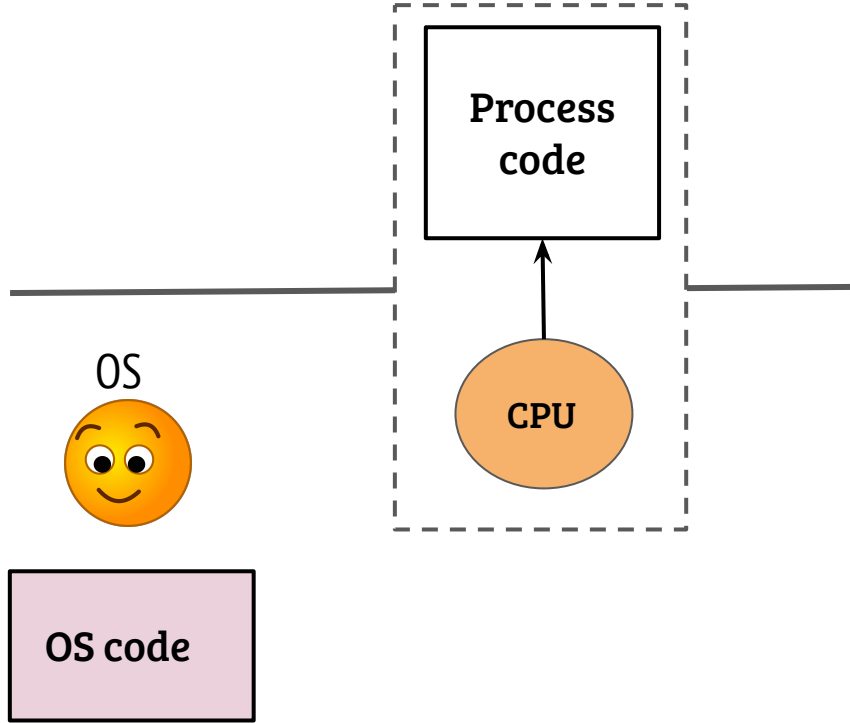
- When we execute “a.out” on a shell a **process control block (PCB)** is created
- Does it raise some questions related to the exact working?

Process creation: What and How?

- How does OS come into action after typing “./a.out” in a shell?




System call




- CPU executing *user code* can invoke the *OS functions* using system calls

System call




- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call

System call



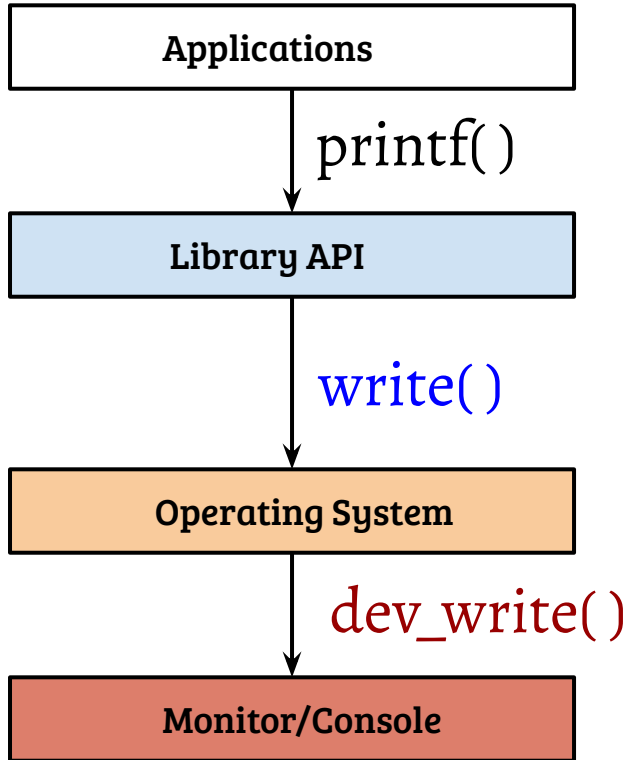
- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call
- How system call is different from a function call?

System call



- CPU executing *user code* can invoke the *OS functions* using system calls
- The CPU executes the OS handler for the system call
- How is system call different from a function call?
- Can be thought as an invocation of privileged functions (will revisit)

System calls and user libraries



- Most system calls are invoked through wrapper library functions
- However, all system calls can be invoked directly
 - For example, in Linux systems, `syscall()` wrapper can be used (Refer: `man syscall`)

A simple system call: getpid()

USER



```
main()  
{  
    printf("%d\n", getpid());  
}
```

pid_t getpid()

{

PCB *current = get_current_process();

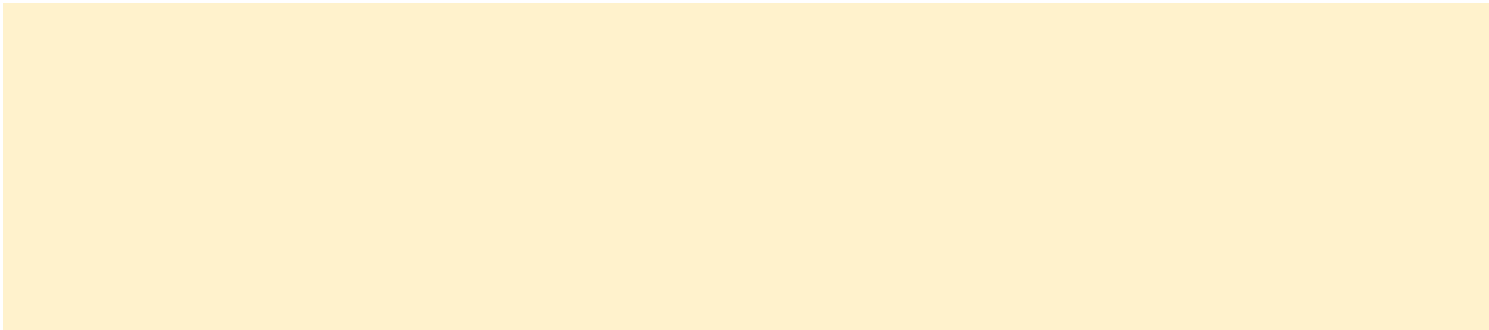
return (current → pid);

}


OS



Process creation: What and How?

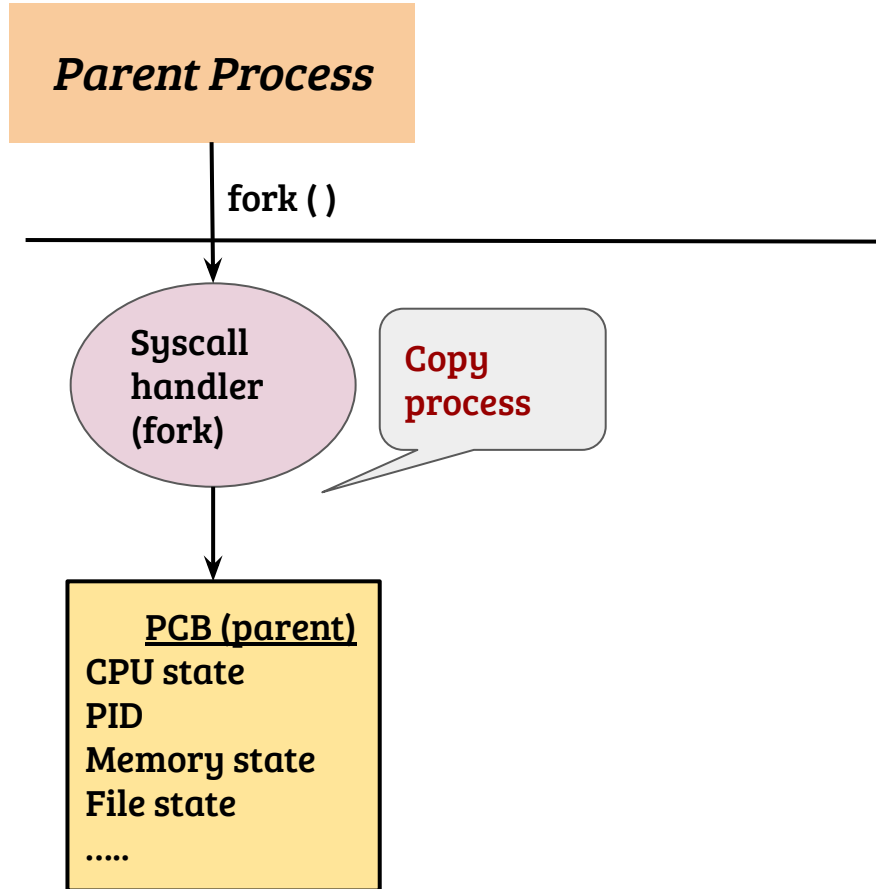
- How does OS come into action after typing “./a.out” in a shell?
 - System calls invoked to explicitly give control to the OS
 - What exact system calls are invoked?
- 

Process creation - fork()

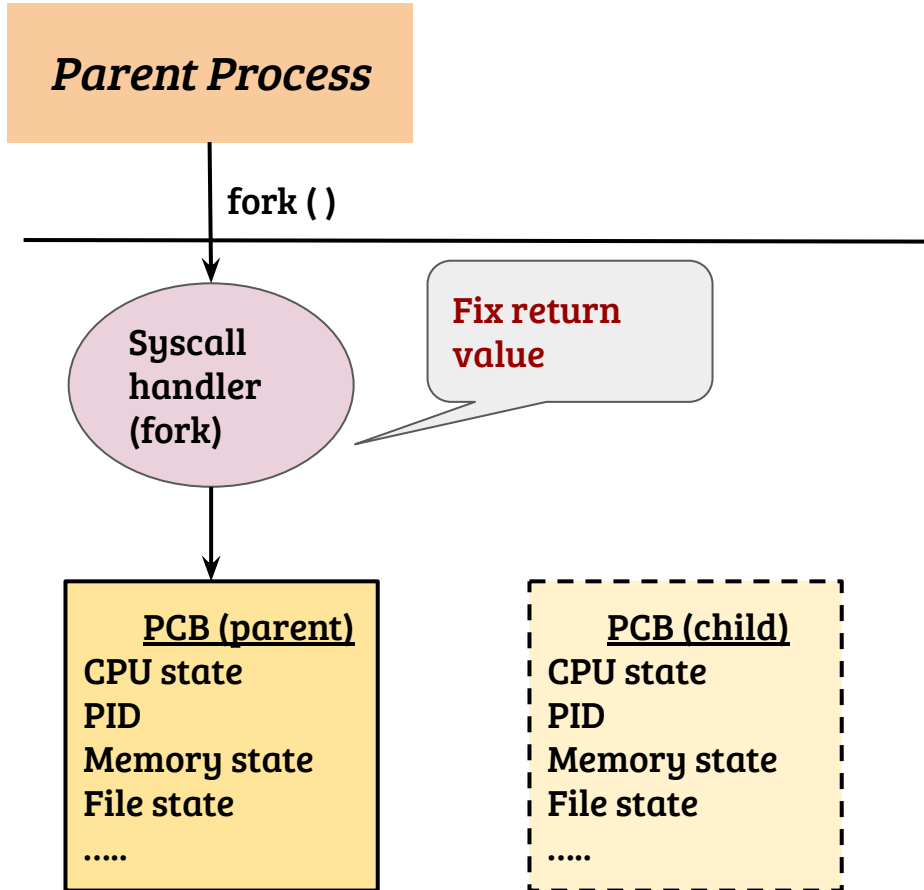


- fork() system call is weird; not a typical “privileged” function call
- fork() creates a new process; a *duplicate* of calling process
- On success, fork
 - Returns PID of child process to the caller (parent)
 - Returns 0 to the child

Typical implementation of fork

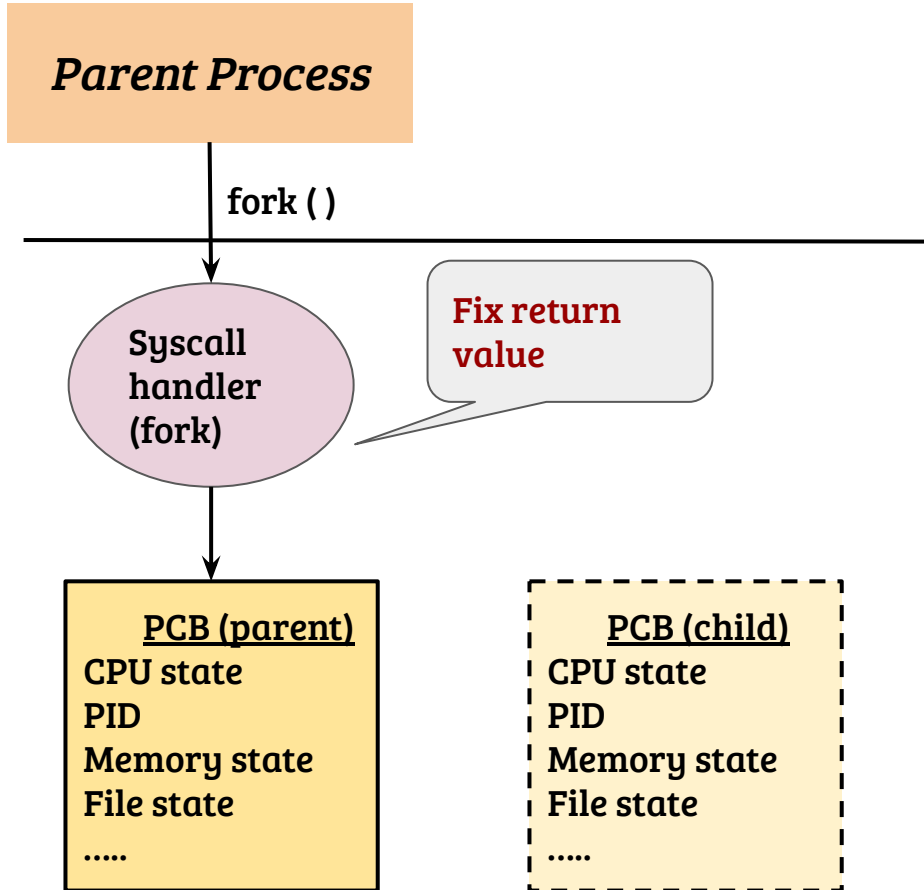


Typical implementation of fork



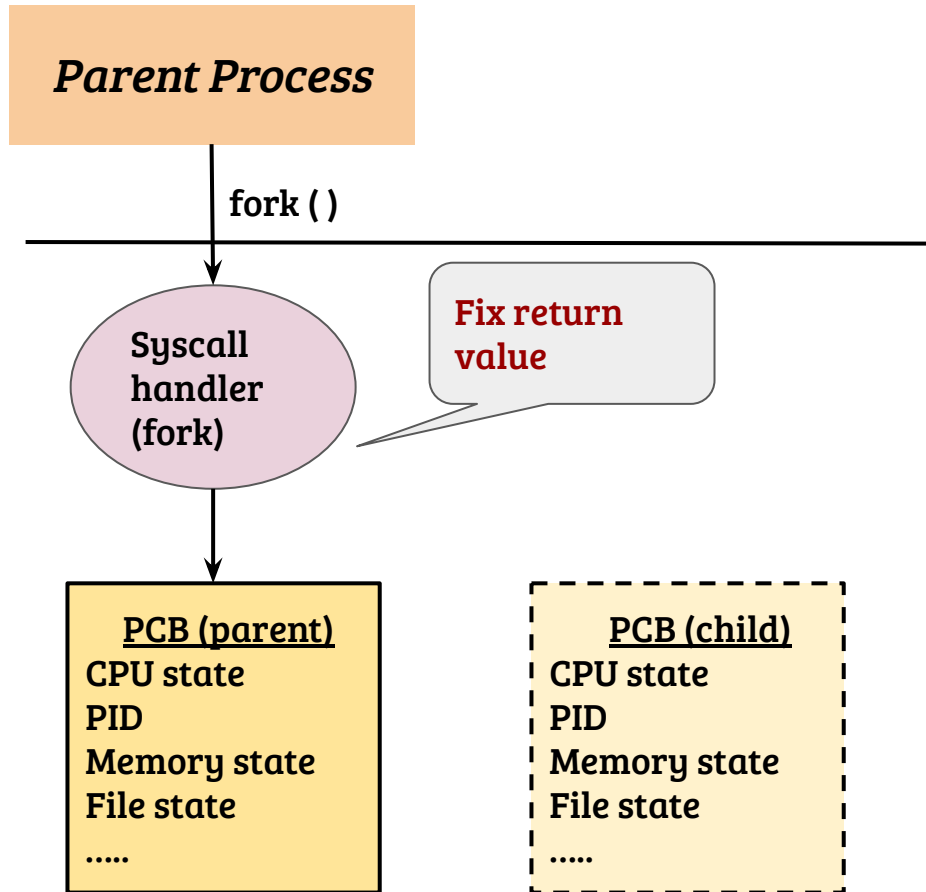
- Child should get '0' and parent gets PID of child as return value. How?

Typical implementation of fork



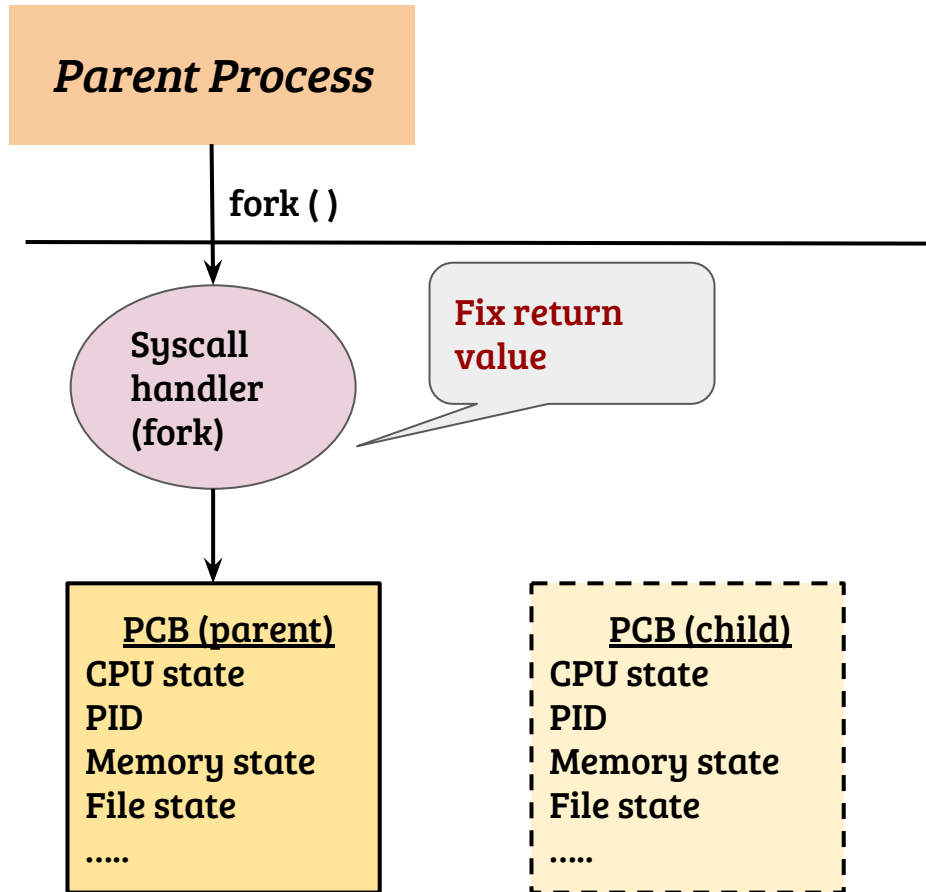
- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child

Typical implementation of fork



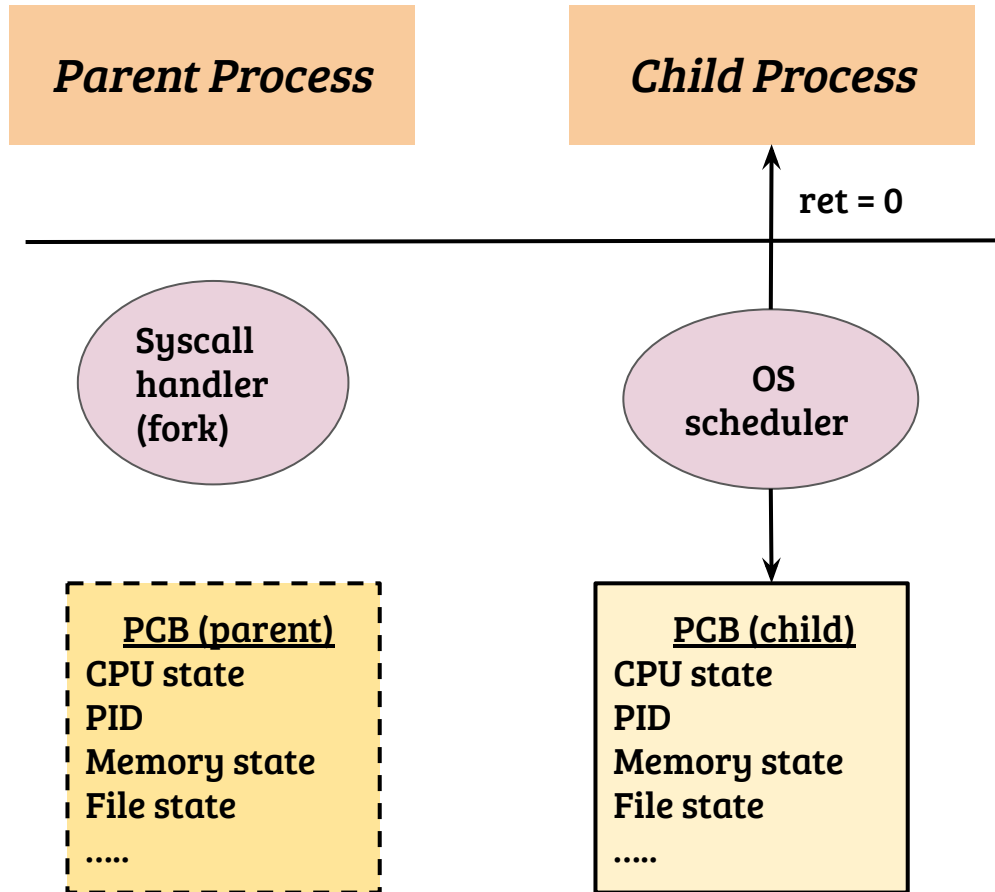
- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?

Typical implementation of fork




- Child should get '0' and parent gets PID of child as return value. How?
- OS returns different values for parent and child
- When does child execute?
- When OS schedules the child process

Typical implementation of fork




- PC is next instruction after `fork()` syscall, for both parent and child
- Child memory is an exact copy of parent
- Parent and child diverge from this point

Load a new binary - `exec()`




- Replace the calling process by a new executable
 - Code, data etc. are replaced by the new process
 - Usually, open files remain open

Typical implementation of exec




- The calling process commits self destruction! (almost)

Typical implementation of exec



- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same

Typical implementation of exec

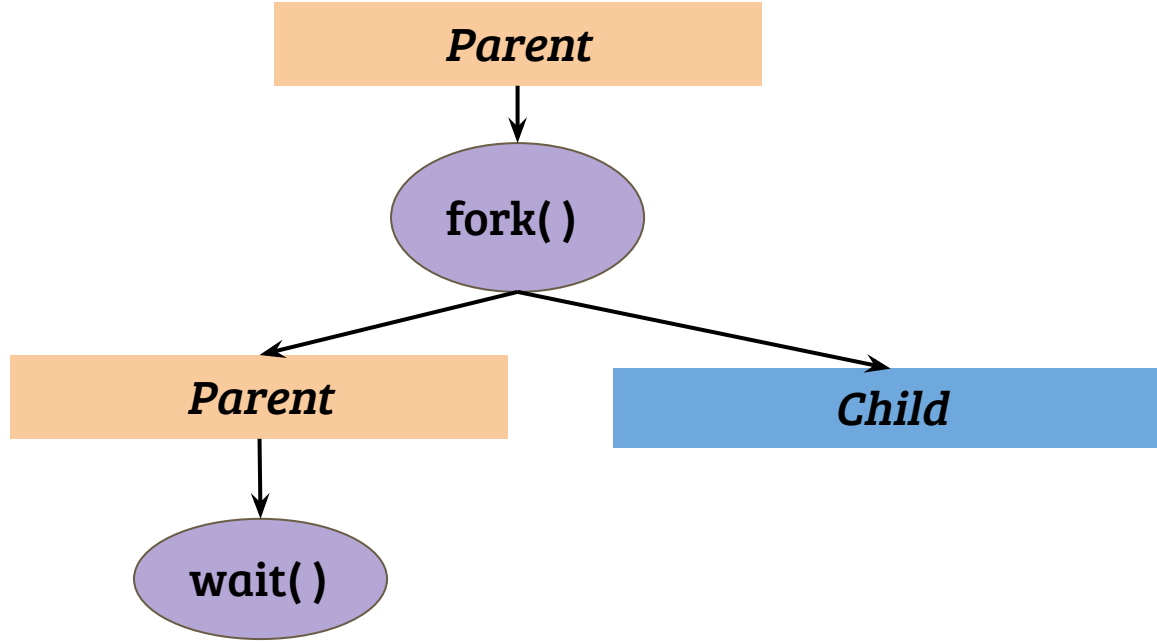


- The calling process commits self destruction! (almost)
- The calling process is cleaned up and replaced by the new executable
- PID remains the same
- On return, new executable starts execution
- PC is loaded with the starting address of the newly loaded binary

Process creation: What and How?

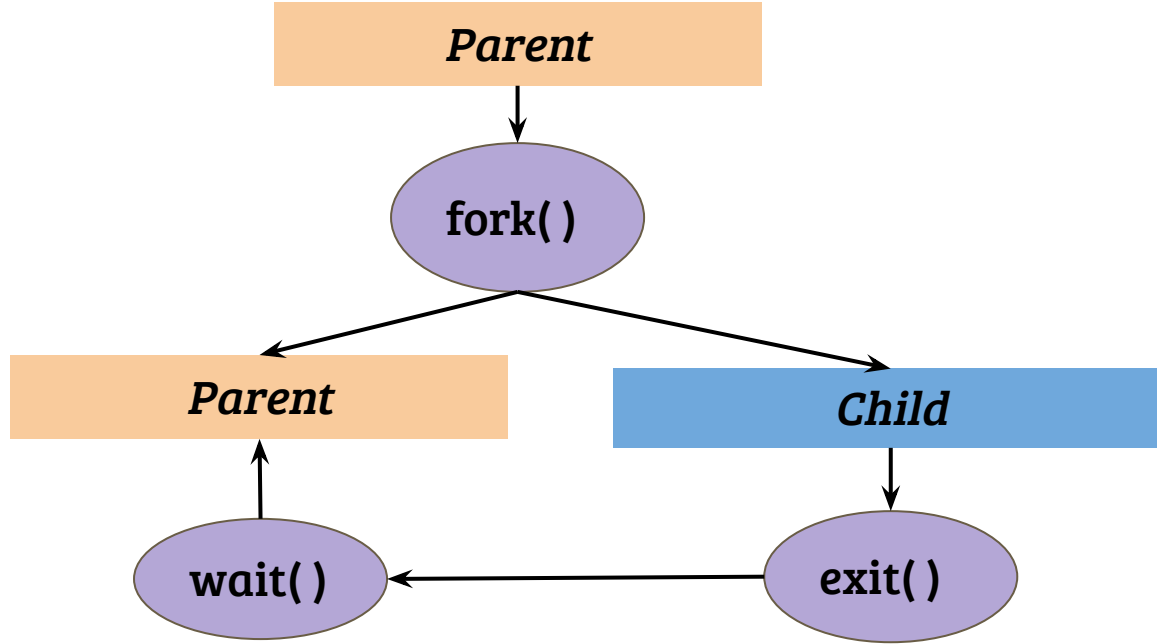
- How does OS come into action after typing “./a.out” in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- `fork()`, `exec ()`, `wait()` and `exit()`
- Who invokes the system calls? In what order?

wait() and exit()




- The wait system call makes the parent wait for child process to exit

wait() and exit()



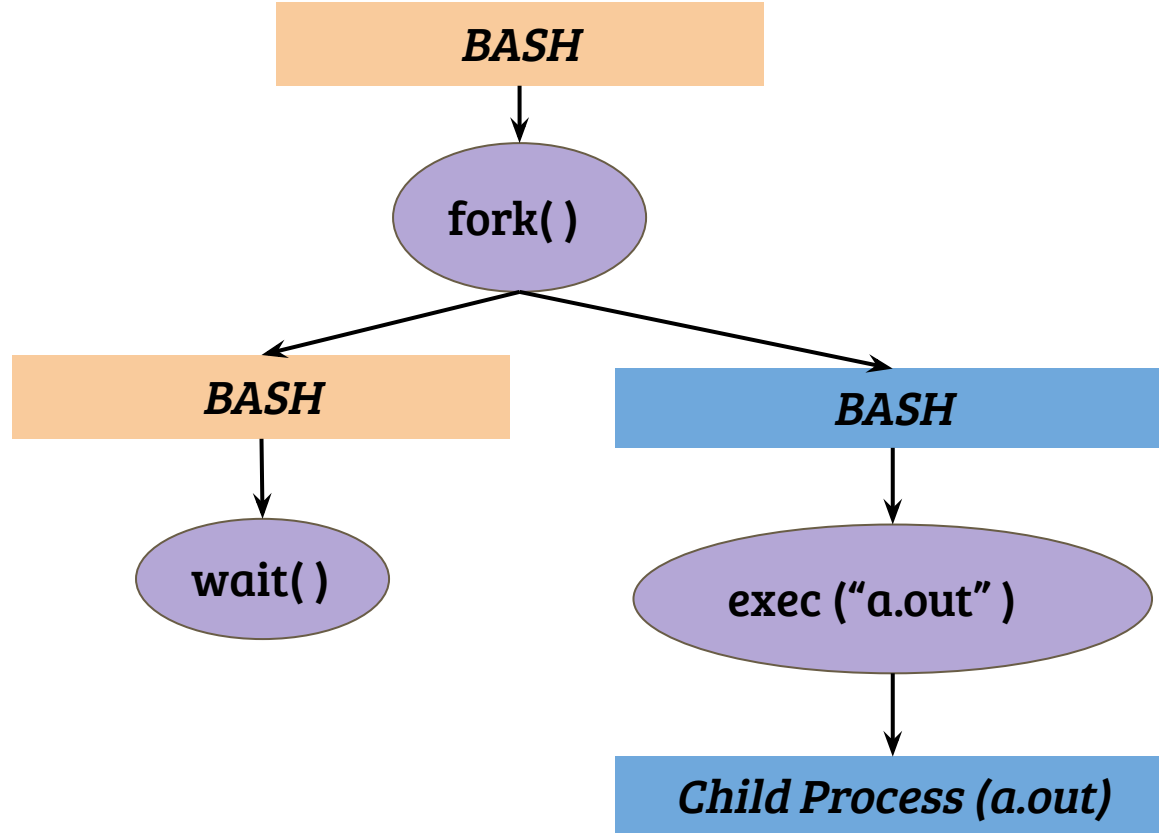
- The wait system call makes the parent wait for child process to exit
- On child **exit()**, the **wait()** system call returns in parent

Shell command line: fork + exec + wait



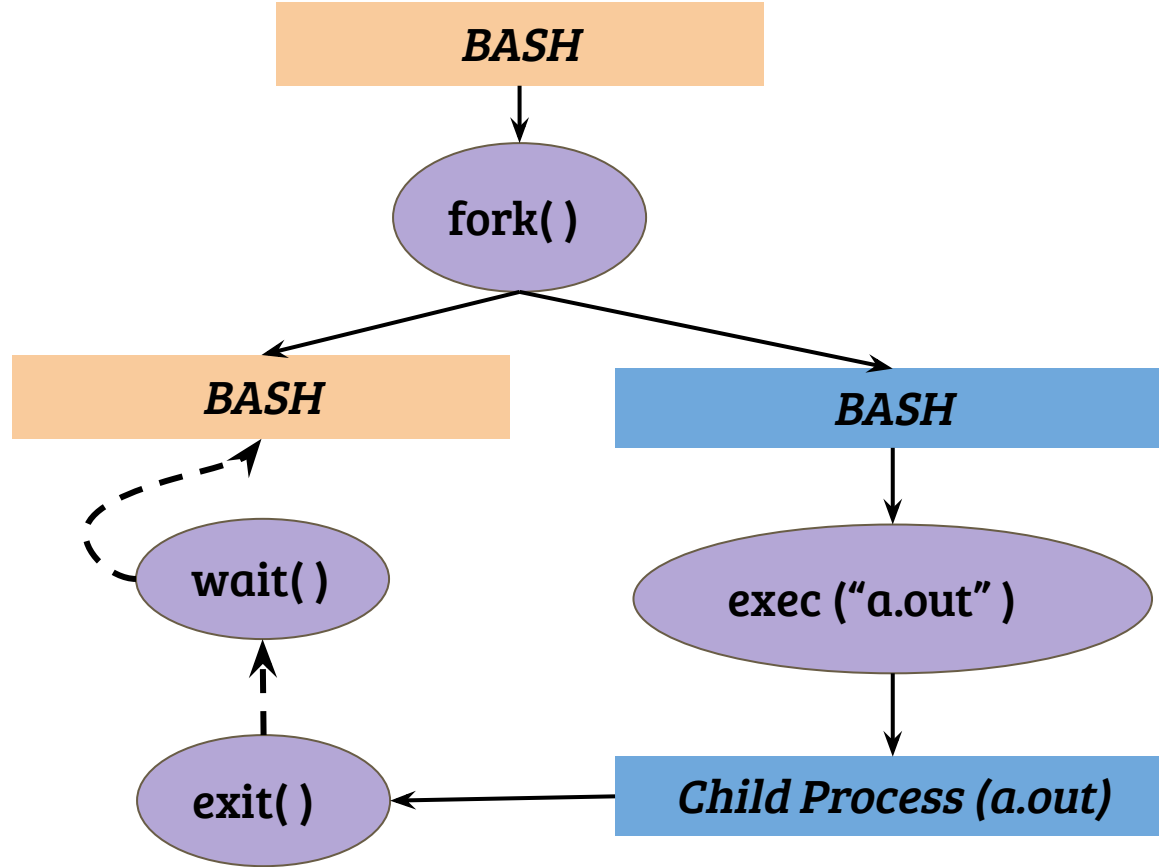
- The BASH process calls *fork()*

Shell command line: fork + exec + wait



- Parent process calls `wait()` to wait for child to finish
- Child process invokes `exec()`

Shell command line: fork + exec + wait




- When child exits, parent gets notified
- The BASH shell is ready for the next command at this point of time

Process creation: What and How?

- How does OS come into action after typing “./a.out” in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- fork(), exec (), wait() and exit()
- Who invokes the system calls? In what order?
- The shell process (bash process)
- What is the first user process?

Unix process family using fork + exec



- Fork and exec are used to create the process tree
- Commands: ps, pstree
- See the /proc directory in linux systems

Process creation: What and How?

- How does OS come into action after typing “./a.out” in a shell?
- System calls invoked to explicitly give control to the OS
- What exact system calls are invoked?
- fork(), exec (), wait() and exit()
- Who invokes the system calls?
- The shell process (bash process)
- What is the first user process?
- In Unix systems, it is called the *init* process
- Who creates and schedules the init process?