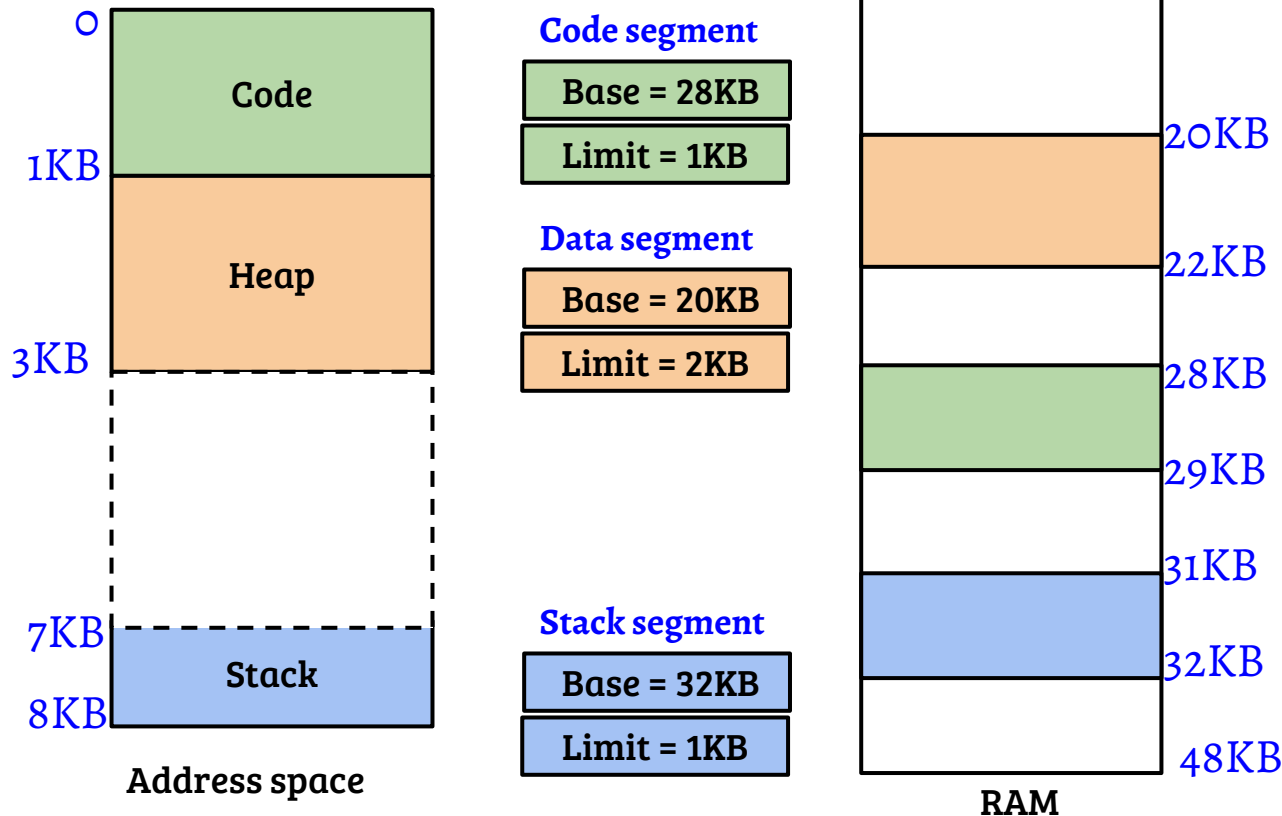


CS330: Operating Systems

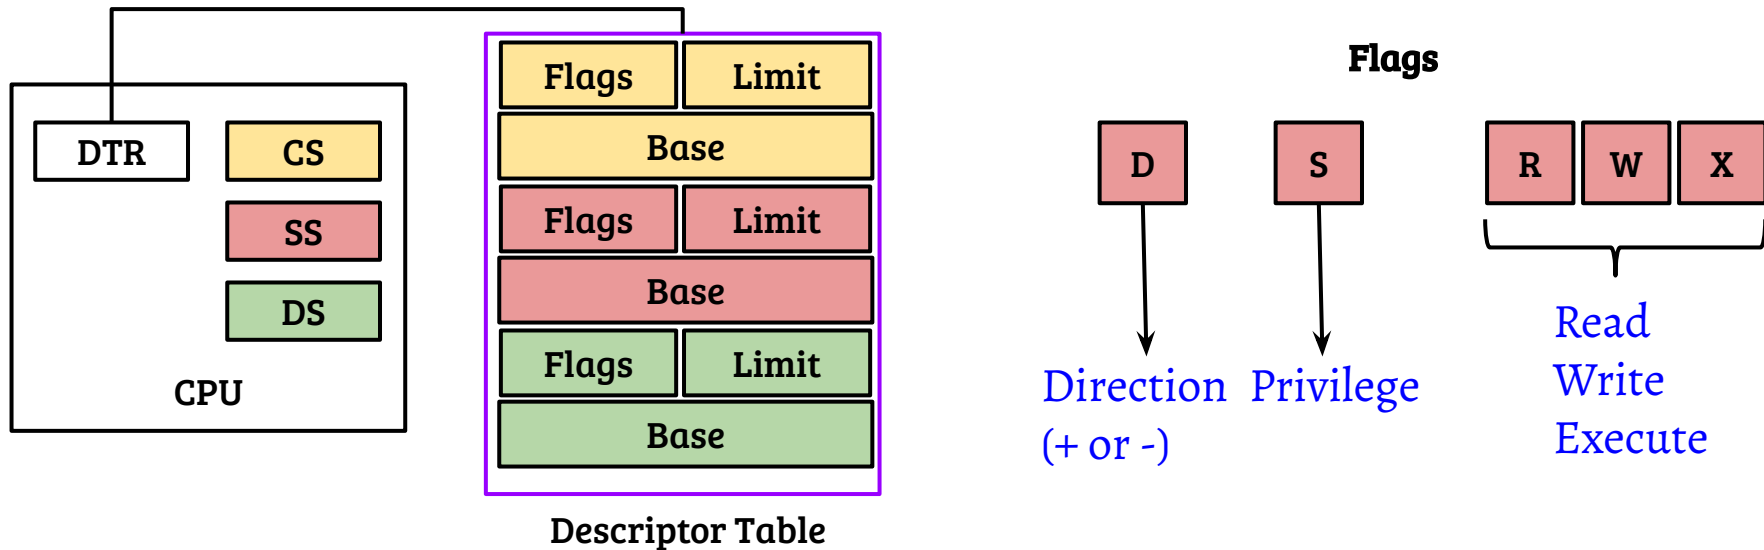
Virtual memory: Paging

Recap: Segmentation



- Extension of the scheme for translation and address space granularity
- Base-limit register pairs per segment

Recap: Segmentation in reality



- Descriptor table register (DTR) is used to access the descriptor table
- # of descriptors depends on architecture
- Separate descriptors used for user and kernel mode

Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it page)
 - Physical memory partitioned in a similar way (call it page frame)


Paging

- Paging addresses the following issues with segmentation
 - External fragmentation caused due to variable sized segments
 - No support for discontinuous/sparse mapping
- The idea of paging
 - Partition the address space into fixed sized blocks (call it pages)
 - Physical memory partitioned in a similar way (call it page frames)
 - OS creates a mapping between *page* to *page frame*
 - H/W uses the mapping to translate VA to PA

Paging example (pages)





- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128




Process address
space

Paging example (pages)



Process address
space


- Virtual address size = 32KB, Page size = 256 bytes
- Address length = 15 bits {0x0 - 0x7FFF}
- # of pages = 128





Virtual address

- Example: For Virtual address *0x0510*, Page number = 5, offset = 16

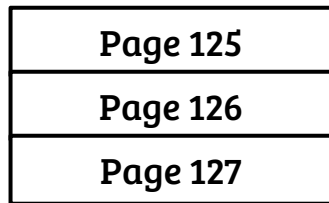
Paging example (page frames)



- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



DRAM



Process address
space

32KB


Paging example (page frames)

0	Page 0
256	Page 1
512	Page 2
768	Page 3
1024	

Page 125
Page 126
Page 127

Process address space

- Physical address size = 64KB
- Address length = 16 bits {0x0 - 0xFFFF}
- # of page frames = 256



Physical address

PFN 0	0
PFN 1	256
PFN 2	512
PFN 3	768
	1024

PFN 253
PFN 254
PFN 255

64KB

DRAM

- Example: For physical address *0x1F51*, PFN = 31, offset = 81

Paging example (page table mapping)

0
256
512
768
1024

Page 0
Page 1
Page 2
Page 3

Page 125
Page 126
Page 127

Process address
space

Page table

1
-
2
4
-

-
3

128 entries

0
256
512
768
1024
1280

PFN 0
PFN 1
PFN 2
PFN 3
PFN 4



PFN 253
PFN 254
PFN 255

DRAM



64KB

- Each entry in page table is called page table entry (PTE)
- Example mapping: page 0 \Rightarrow PFN 1, page 2 \Rightarrow PFN 2 and so on


Paging example (page table walk)



```
PTW (vaddr V, PTable P)
// Input: Virtual address, Page table
// Returns physical address
{
    Entry = P[V >> 8];
    if (Entry.present)
        return (Entry.PFN << 8) + (V & 0xFF);
    Raise PageFault;
}
```

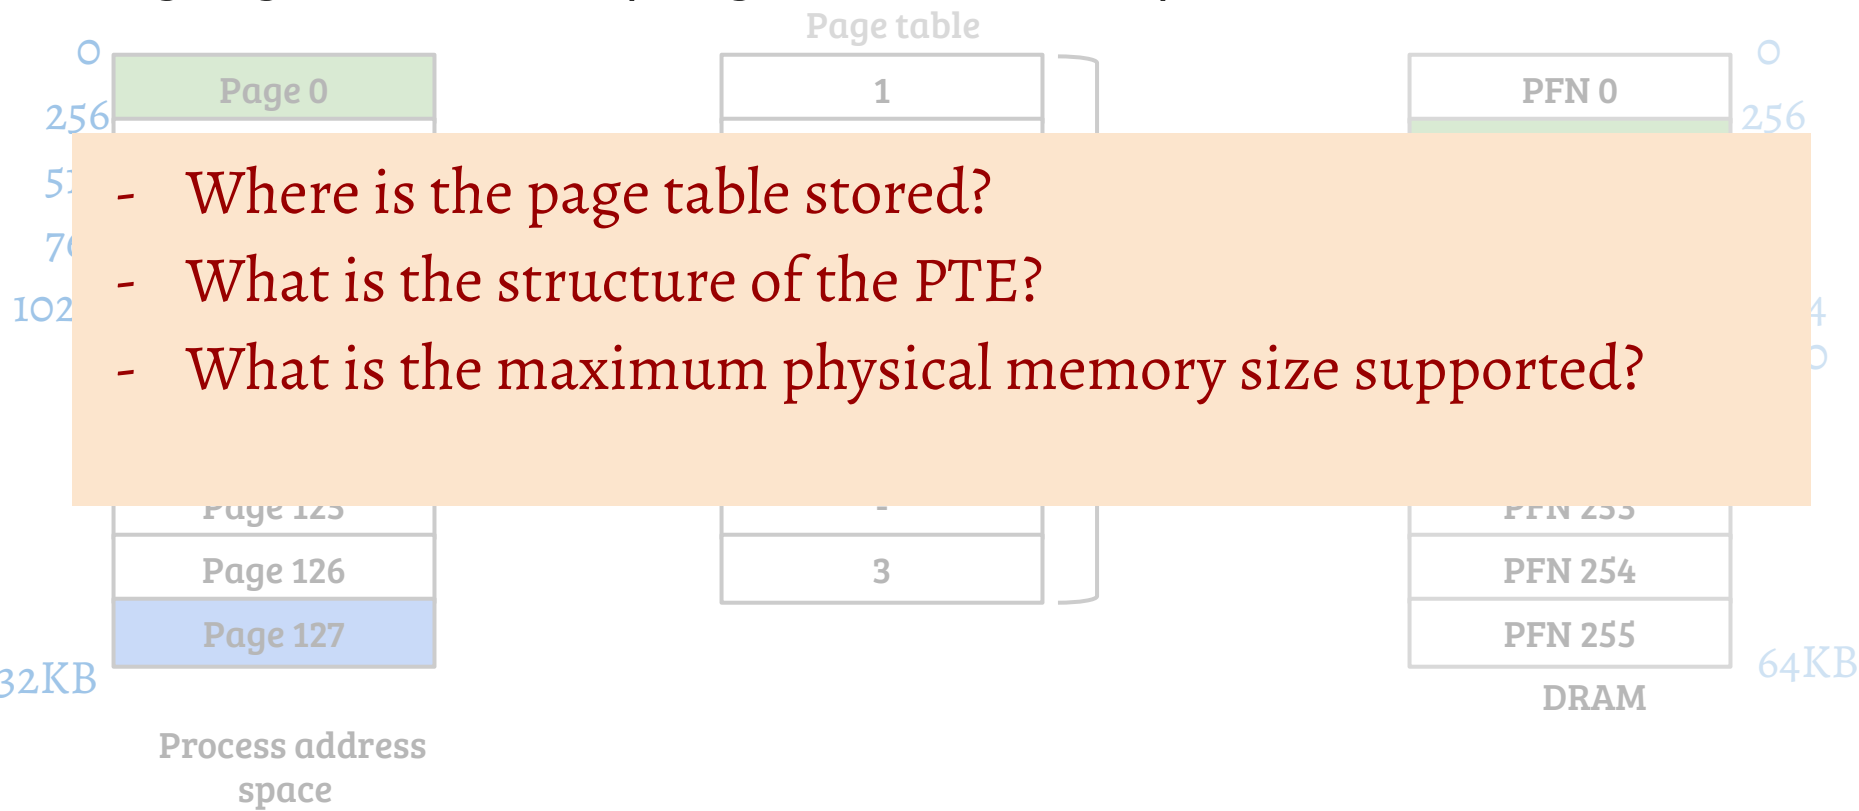


Paging example (example translation)

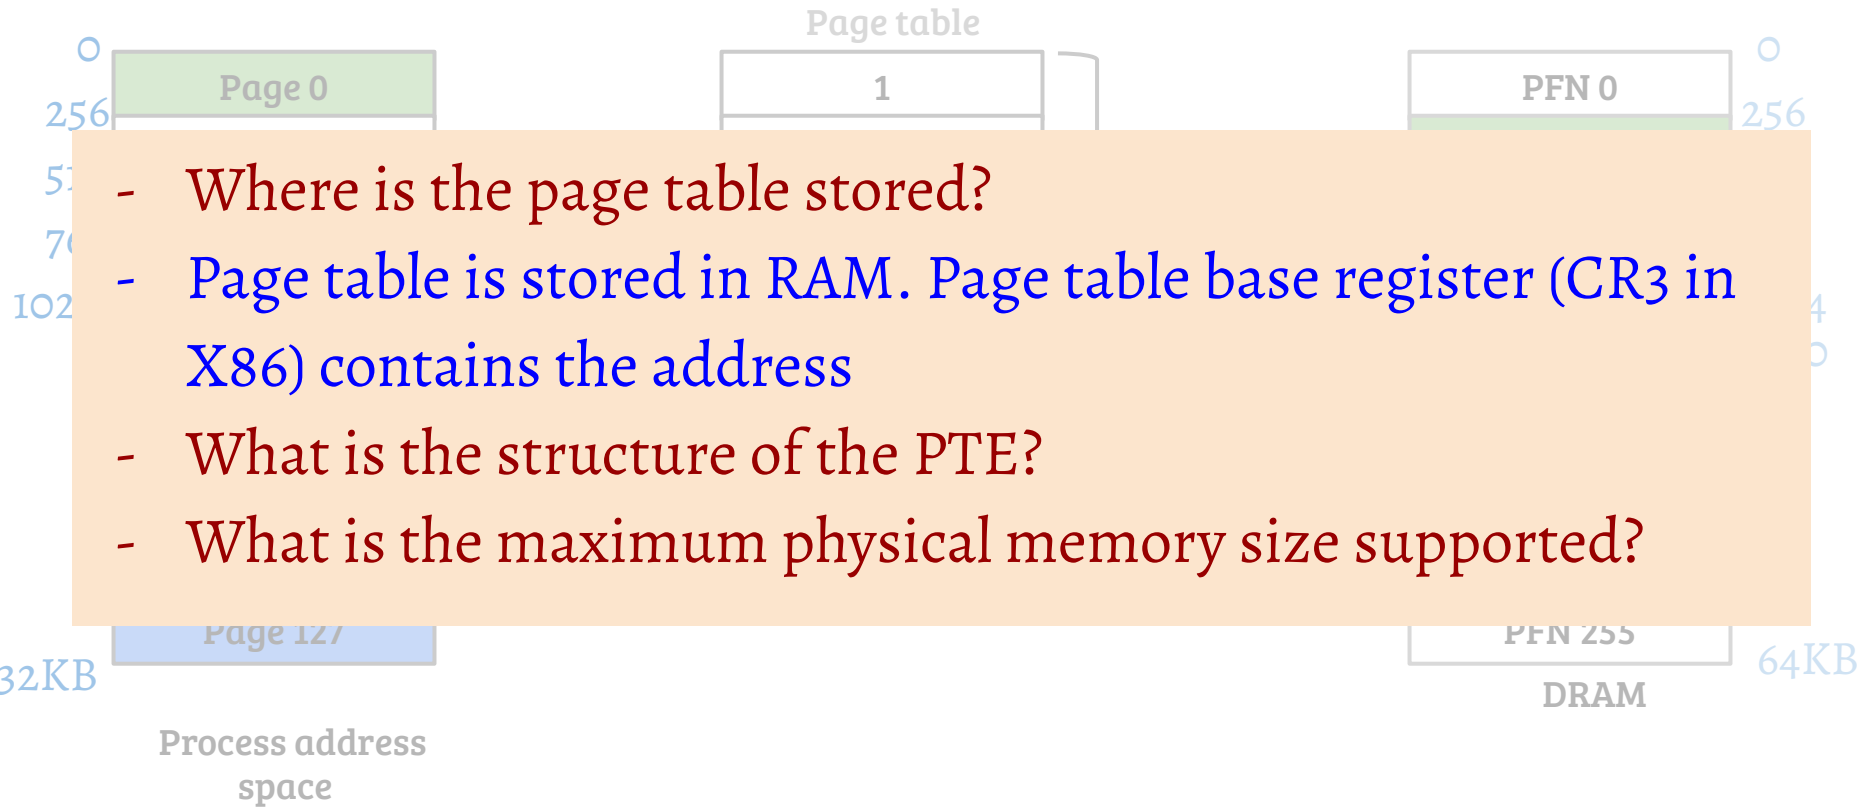


- Process address space
- Virtual address 0x10 translates to physical address 0x110
 - Virtual address 0x7FF0 translates to physical address 0x3F0

Paging example (page table walk)



Paging example (page table walk)



Paging example (structure of an example PTE)



- PFN occupies a significant portion of PTE entry (8 bits in this example)

P

Present bit, 1 \Rightarrow entry is valid

W

Write bit, 1 \Rightarrow Write allowed

S

Privilege bit, 0 \Rightarrow only kernel mode access is allowed

A

Accessed bit, 1 \Rightarrow Address accessed (set by H/W during walk)

D

Dirty bit, 1 \Rightarrow Address written (set by H/W during walk)

X

Execute bit, 1 \Rightarrow Instruction fetch allowed for this page

Reserved/unused bits

Paging example (Page table entries)

0
256
512
768
1024

Page 0
Page 1
Page 2
Page 3

Page table

0x125
0x0
0x207
0x407
0x0

0
256
512
768
1024
1280

PFN 0
PFN 1
PFN 2
PFN 3
PFN 4

Page 125
Page 126
Page 127

0x0
0x307

PFN 253
PFN 254
PFN 255

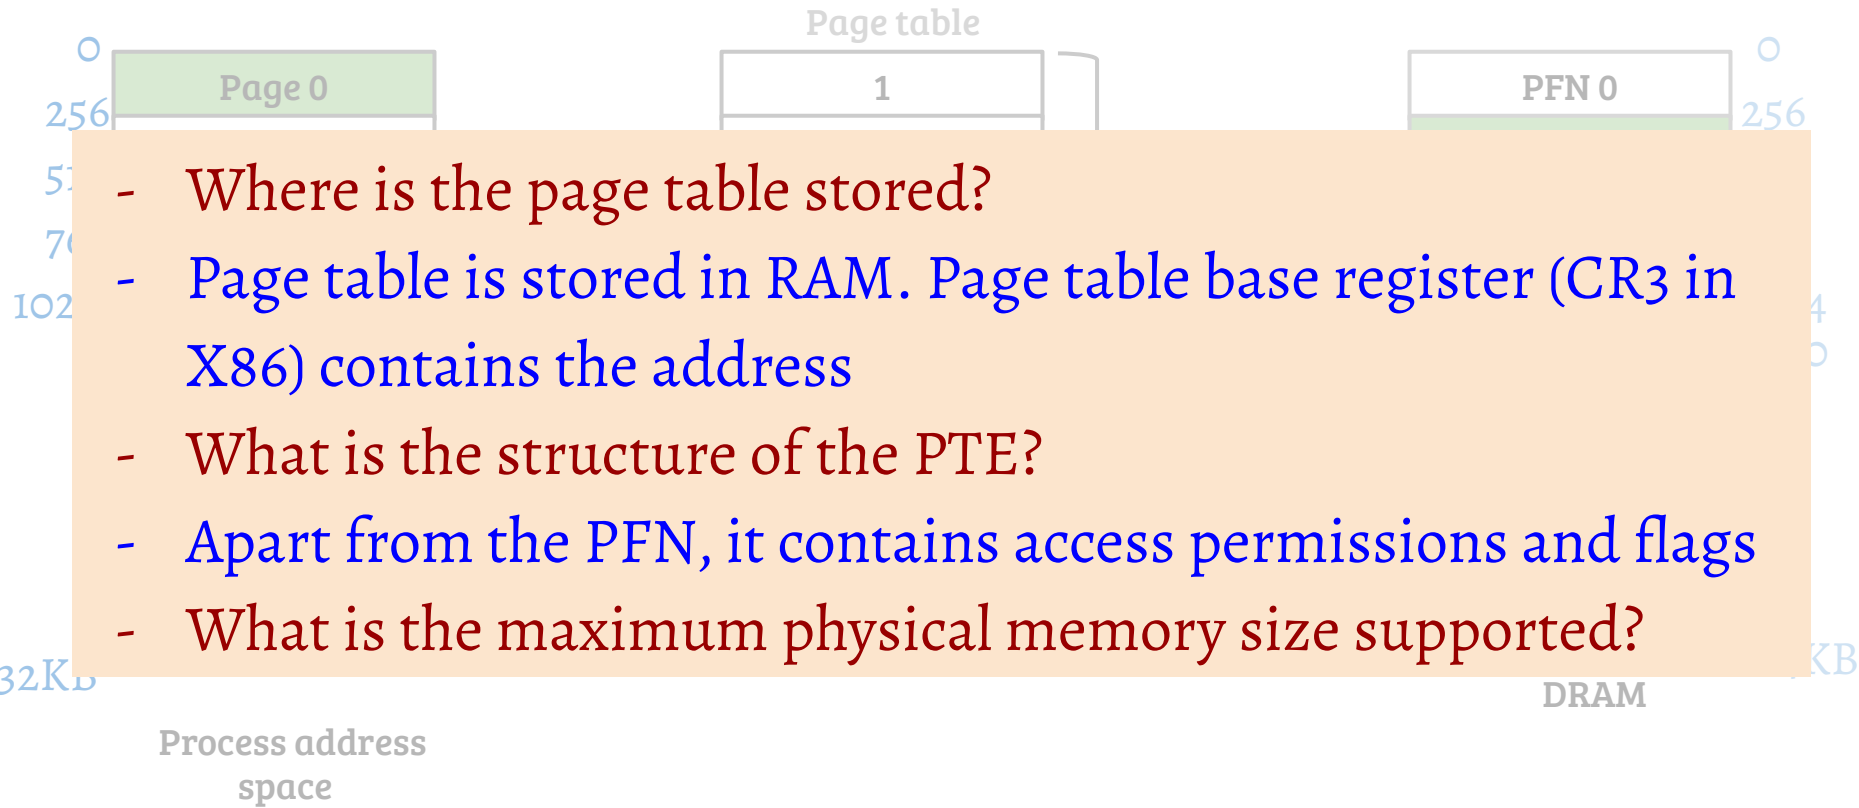
64KB

DRAM

Process address
space

- Code: Page 0 (Read and Execute)
- Data: Page 2 and Page 3 (Read and Write)
- Stack: Page 127 (Read and Write)

Paging example (page table walk)



Paging example (page table walk)

- Where is the page table stored?
- Page table is stored in RAM. Page table base register (CR3 in X86) contains the address
- What is the structure of the PTE?
- Apart from the PFN, it contains access permissions and flags
- What is the maximum physical memory size supported?
- For this example, 8-bits can be used to specify 256 page frames. Maximum RAM size = $256 * 256 = 64\text{KB}$

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?

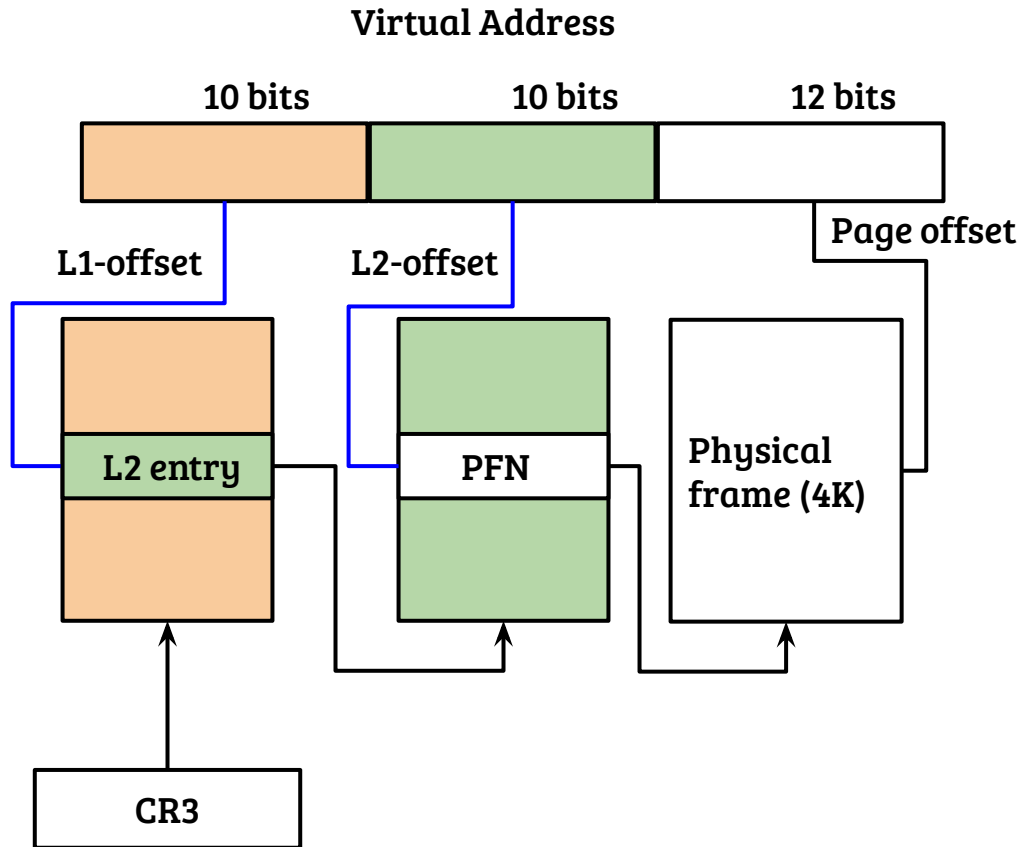
Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system?

Paging: one level of page table may not be feasible!

- Consider a 32-bit address space (=4GB)
- What should be the page size for this system?
- Large page size results in *internal fragmentation*
- Assuming page size = 4KB, How many entries are required in a one-level paging system? (2^{20} entries)
- Not possible to hold 2^{20} entries in a single page
- Therefore, multi-level page tables are used in modern systems

Two-level page tables (32-bit virtual address)



- Two-level page table
- Level-1 page table contains entries pointing to Level-2 page table structures
- Level-2 entry contains PFN along with flags