CS330: Operating Systems

Locking

Recap: Synchronization and locking

- Locking is necessary when multiple contexts access shared resources
- Example: Multiple threads, multiple OS execution contexts
- Efficiency of lock and unlock operations
- Hardware-assisted lock implementations are used for efficiency
- Lock acquisition delay vs. wasted CPU cycles
- Use waiting locks and spinlocks depending on the requirement
- Fairness of the locking scheme
- Contending threads should not starve for the lock (infinitely)

Recap: Synchronization and locking

- Locking is necessary when multiple contexts access shared resources
- Example: Multiple threads, multiple OS execution contexts
- Efficiency of lock and unlock operations
- Hardware-assisted lock implementations are used for efficiency
- Lock acquisition delay vs. wasted CPU cycles
- Use waiting locks and spinlocks depending on the requirement
- Fairness of the locking scheme
- Contending threads should not starve for the lock

Agenda: Spinlocks, Semaphore and mutex (waiting locks)

Spinlock: Buggy attempt

1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.   while(*L);
5.   *L = 1;
6. }
7. unlock(L)
8. {
9.   *L = 0;
10. }

- Does this implementation work?

Spinlock: Buggy attempt

1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.   while(*L);
5.   *L = 1;
6. }
7. unlock(L)
8. {
9.   *L = 0;
10. }

- Does this implementation work?
- No, it does not ensure *mutual exclusion*
- Why?

Spinlock: Buggy attempt

1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.   while(*L);
5.   *L = 1;
6. }
7. unlock(L)
8. {
9.   *L = 0;
10. }

- Does this implementation work?
- No, it does not ensure *mutual exclusion*
- Why?
  - Single core: Context switch between line #4 and line #5
  - Multicore: Two cores exiting the while loop by reading lock = 0

Spinlock: Buggy attempt

1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.   while(*L);
5.   *L = 1;
6. }
7. unlock(L)
8. {
9.   *L = 0;
10. }

- Does this implementation work?
- No, it does not ensure *mutual exclusion*
- Why?
  - Single core: Context switch between line #4 and line #5
  - Multicore: Two cores exiting the while loop by reading lock = 0
- Core issue: Compare and swap has to happen atomically!

Spinlock using atomic exchange

```
1.  lock_t *L; // Initial value = 0
2.  lock(L)
3.  {
4.      while(atomic_xchg(*L, 1));
5.  }
6.  unlock(L)
7.  {
8.      *lock = 0;
9.  }
```

- Atomic exchange: exchange the value of memory and register atomically
- atomic_xchg (int *PTR, int val) returns the value at PTR before exchange
- Ensures mutual exclusion if "val" is stored on a register
- No fairness guarantees

Spinlock using XCHG on X86

```
lock(lock_t *L )
{
    asm volatile(
    "mov $1, %%rax;"
    "loop: xchg %%rax, (%%rdi);"
    "cmp $0, %%rax;"
     "jne loop;"
      : : : "memory" );
}
unlock(int *L ) { *L = 0;}
```

- XCHG R, M ⇒ Exchange value of register R and value at memory address M

- RDI register contains the lock argument

- Exercise: Visualize a context switch between any two instructions and analyse the correctness

Spinlock using compare and swap

1. lock_t *L; // Initial value = 0
2. lock(L)
3. {
4.   while(CAS(*L, 0, 1));
5. }
6. unlock(L)
7. {
8.   *lock = 0;
9. }

- Atomic compare and swap: perform the condition check and swap atomically
- CAS (int *PTR, int cmpval, int newval) sets the value of PTR to newval if cmpval is equal to value at PTR. Returns 0 on successful exchange
- No fairness guarantees!

CAS on X86: cmpxchg

cmpxchg source[Reg] destination [Mem/Reg]
Implicit registers : rax and flags

1.   if rax == [destination]
2.   then
3.   flags[ZF] = 1
4.   [destination] = source
5.   else
6.   flags[ZF] = 0
7.   rax = [destination]

- "cmpxchg" is not atomic in X86, should be used with a "lock" prefix

Spinlock using CMPXCHG on X86

```
lock(lock_t *L )
{
asm volatile(
    "mov $1, %%rcx;"
    "loop: xor %%rax, %%rax;"
    "lock cmpxchg %%rcx, (%%rdi);"
    "jnz loop;"
    : : : "rcx", "rax", "memory");
}
unlock(lock_t *L ) { *L = 0;}
```

- Value of RAX (=0) is compared against value at address in register RDI and exchanged with RCX (=1), if they are equal
- Exercise: Visualize a context switch between any two instructions and analyse the correctness

Load Linked (LL) and Store conditional (SC)

- LoadLinked (R, M)
    - Like a normal load, it loads R with value of M
    - Additionally, the hardware keeps track of future stores to M
- StoreConditional (R, M)
    - Stores the value of R to M if no stores happened to M after the execution of LL instruction (after execution, R = 1)
    - Otherwise, store is not performed (after execution R=0)
- Supported in RISC architectures like mips, risc-v etc.

Spinlock using LL and LC

```
lock_t *L;    //initial value = 0
lock( lock_t *L)
{
 while(LoadLinked(L) ||
     !StoreConditional(L, 1));
}
unlock( lock_t *L) { *L = 0;}
```

```
lock:    LL  R1,  (R2);  //R2 = lock address
         BNEQZ  R1,  lock;
         ADDUI  R1,  R0, #1;    //R1 = 1
         SC R1, (R2)
         BEQZ R1, lock
```

- Efficient as the hardware avoids memory traffic for unsuccessful lock acquire attempts
- Context switch between LL and SC results in SC to fail

Spinlocks: reducing wasted cycles

- Spinning for locks can introduce significant CPU overheads and increase energy consumption
- How to reduce spinning in spinlocks?

Spinlocks: reducing wasted cycles

- Spinning for locks can introduce significant CPU overheads and increase energy consumption
- How to reduce spinning in spinlocks?
- Strategy: Back-off after every failure, exponential back-off used mostly

```
lock( lock_t *L) {
        u64 backoff = 0;
        while(LoadLinked(L) || !StoreConditional(L, 1)){
                if(backoff < 63)  ++backoff;
                        pause(1 << backoff);  // Hint to processor
}
```

Fairness in spinlocks

- Spinlock implementations discussed so far are not fair,
  - no bounded waiting
- To ensure fairness, some notion of ordering is required
- What if the threads are granted the lock in the order of their arrival to the lock contention loop?
  - A single lock variable may not be sufficient
  - Example solution: Ticket spinlocks

Atomic fetch and add (xadd on X86)

xadd   R,   M

$\mathrm{TmpReg}\ T = R + [M]$

$R = [M]$

$[M] = T$

- Example: M = 100; RAX = 200
- After executing "lock xadd %RAX, M", value of RAX = 100, M = 300
- Require lock prefix to be atomic

Ticket spinlocks (OSTEP Fig. 28.7)

```
struct lock_t{
        long ticket;
        long turn;
};
void init_lock (struct lock_t *L){
    L → ticket = 0;  L → turn = 0;
}
void unlock(struct lock_t *L){
        L → turn++;
}
```

```
void lock(struct lock_t *L){
    long myturn = xadd(&L → ticket, 1);
    while(myturn != L → turn)
        pause(myturn - L → turn);
}
```
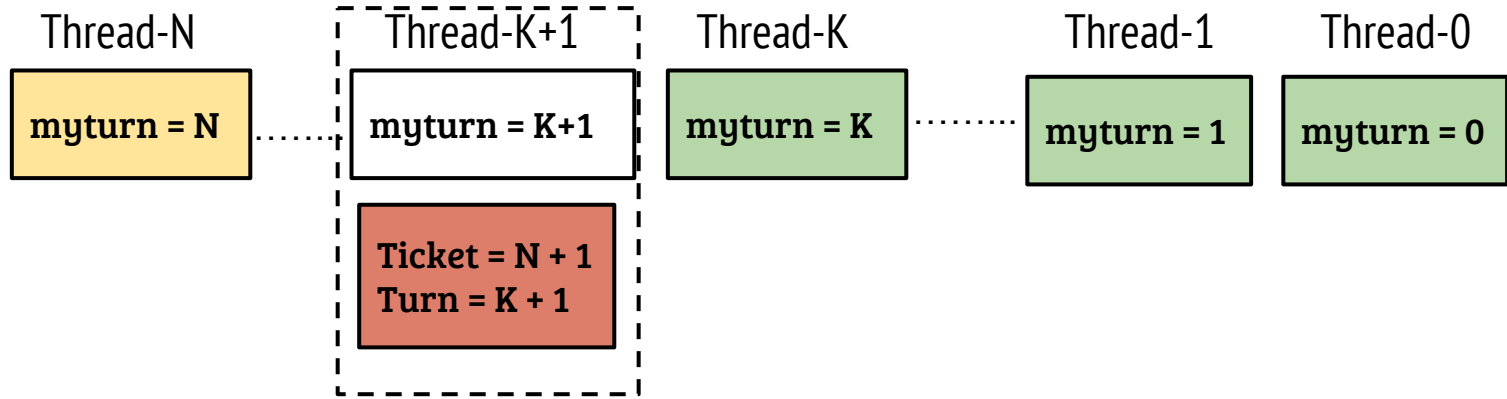
- Example: Order of arrival: T1 T2 T3
- T1 (in CS) : myturn = 0, L = {1, 0}
- T2: myturn = 1, L = {2, 0}
- T3: myturn = 2, L = {3,0}
- T1 unlocks, L = {3, 1}. T2 enters CS

Ticket spinlock

- Local variable "myturn" is equivalent to the order of arrival
- If a thread is in CS ⟹ Local Turn must be same as "Turn"
- Threads waiting = Ticket - Turn -1

Ticket spinlock

Thread-N	Thread-K+1	Thread-K	Thread-1	Thread-0
myturn = N	myturn = K+1	myturn = K	myturn = 1	myturn = 0
	Ticket = N + 1 Turn = K + 1			

- Value of turn incremented on lock release
- Thread which arrived just after the current thread enters the CS
- When a new thread arrives, it gets the lock after the other threads ahead of the new thread acquire and release the lock

Ticket spinlock

Thread-N+1	Thread-N	Thread-K+1	Thread-K	Thread-1	Thread-0
myturn = N+1	myturn = N	myturn = K+1	myturn = K	myturn = 1	myturn = 0

Ticket = N + 2
Turn = K + 1

- Ticket spinlock guarantees bounded waiting
- If N threads are contending for the lock and execution of the CS consumes T cycles, then bound = N * T (assuming negligible context switch overhead)

Ticket spinlock (with yield)

```
void lock(struct lock_t *L){
    long myturn = xadd(&L → ticket, 1);
    while(myturn != L → turn)
        sched_yield( );
}
```

- Why spin if the thread's turn is yet to come?
- Yield the CPU and allow the thread with ticket (or other non contending threads)
- Further optimization
  - Allow the thread with "myturn" value one more than "L→ turn" to continue spinning