

# Computer Networks I


## Network Layer Details - 2

Amitangshu Pal  
Computer Science and Engineering  
IIT Kanpur


# Routing protocols

**Routing protocol goal:** determine "good" paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- **path:** sequence of routers packets traverse from given initial source host to final destination host
- **"good":** least "cost", "fastest", "least congested"



# Graph abstraction: link costs



$c_{a,b}$ : cost of **direct** link connecting a and b  
e.g.,  $c_{w,z} = 5$ ,  $c_{u,z} = \infty$

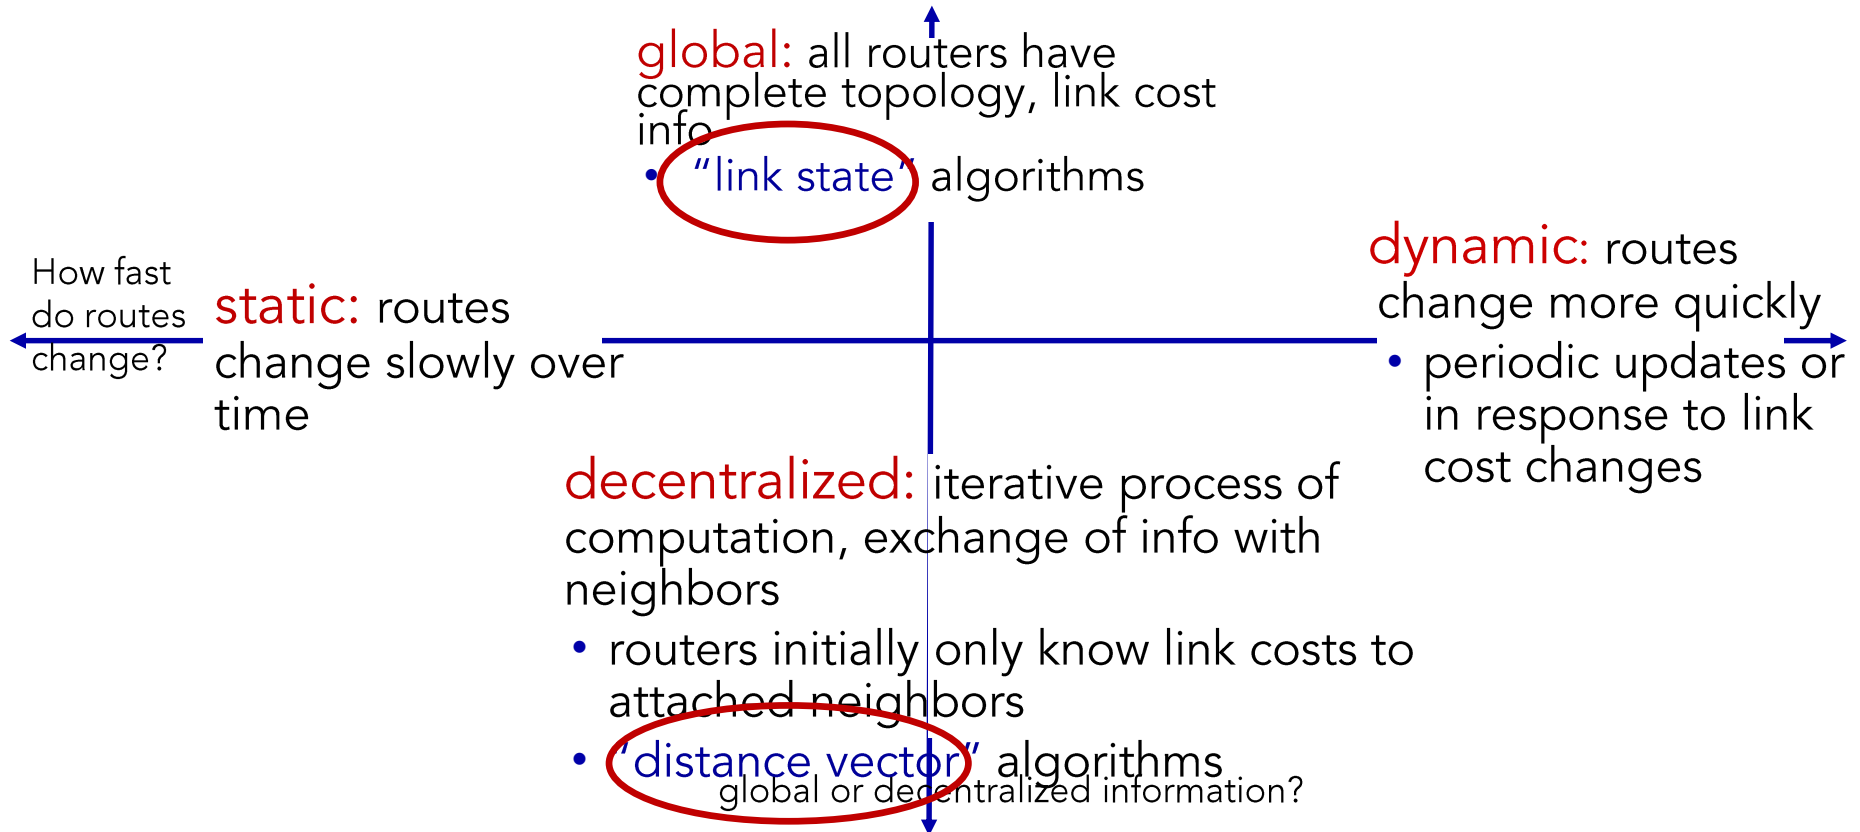
cost defined by network operator:  
could always be 1, or inversely  
related to bandwidth, or inversely  
related to congestion

graph:  $G = (N, E)$

N: set of routers =  $\{ u, v, w, x, y, z \}$

E: set of links =  $\{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$

# Routing algorithm classification



# Link State Routing Protocol

---

# Dijkstra's link-state routing algorithm

- **centralized**: network topology, link costs known to all nodes
  - accomplished via "link state broadcast"
  - all nodes have same info
- computes least cost paths from one node ("source") to all other nodes
  - gives **forwarding table** for that node
- **iterative**: after  $k$  iterations, know least cost path to  $k$  destinations

## notation

- $c_{x,y}$ : direct link cost from node  $x$  to  $y$ ;  $= \infty$  if not direct neighbors
- $D(v)$ : current estimate of cost of least-cost-path from source to destination  $v$
- $p(v)$ : predecessor node along path from source to  $v$
- $N'$ : set of nodes whose least-cost-path definitively known


# Dijkstra's link-state routing algorithm

1 Initialization:

```
2  N' = {u}                                /* compute least cost path from u to all other
nodes */
3  for all nodes v
4    if v adjacent to u                    /* u initially knows direct-path-cost only to direct
neighbors */
5      then D(v) =  $c_{u,v}$                 /* but may not be minimum cost!
*/
6    else D(v) =  $\infty$ 
7
```


8 Loop

```
9  find w not in N' such that D(w) is a minimum
10 add w to N'
11 update D(v) for all v adjacent to w and not in N' :
12    $D(v) = \min ( D(v), D(w) + c_{w,v} )$ 
13 /* new least-path-cost to v is either old least-cost-path to v or known
14 least-cost-path to w plus direct-cost from w to v */
15 until all nodes in N'
```



# Dijkstra's algorithm: an example

Step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2, u	5, u	1, u	$\infty$	$\infty$
1	u, x	2, u	4, x	1, u	2, x	$\infty$
2	u, x, y	2, u	3, y	1, u	2, x	4, y
3	u, x, y, v	2, u	3, y	1, u	2, x	4, y
4	u, x, y, w	2, u	3, y	1, u	2, x	4, y
5	u, x, y, w, z	2, u	3, y	1, u	2, x	4, y




Initialization (step 0): For all  $a$ : if  $a$  adjacent to  $u$  then  $D(a) = c_{u,a}$


find  $a$  not in  $N'$  such that  $D(a)$  is a minimum  
 add  $a$  to  $N'$   
 update  $D(b)$  for all  $b$  adjacent to  $a$  and not in  $N'$  :  
 $D(b) = \min ( D(b), D(a) + c_{a,b} )$



# Dijkstra's algorithm: an example



resulting least-cost-path tree from u:      resulting forwarding table in u:




destination	outgoing link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
x	(u,x)

route from u to v directly

route from u to all other destinations via x

# Dijkstra's algorithm: another example

Step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	7, u	3, u	5, u	$\infty$	$\infty$
1	uw	6, w		5, u	11, w	$\infty$
2	uwvx	6, w			11, w	14, x
3	uwvxv				10, v	14, x
4	uwxvy					12, y
5	uwxvyz					



## Notes:

- Construct least-cost-path tree by tracing predecessor nodes
- Ties can exist (can be broken arbitrarily)

# Dijkstra's algorithm: discussion

Algorithm complexity:  $n$  nodes


- each of  $n$  iteration: need to check all nodes,  $w$ , not in  $N$
- $n(n+1)/2$  comparisons:  $O(n^2)$  complexity
- more efficient implementations possible:  $O(n \log n)$

message complexity:


- each router must **broadcast** its link state information to other  $n$  routers
  - efficient (and interesting!) broadcast algorithms:  $O(n)$  link crossings to disseminate a broadcast message from one source
  - each router's message crosses  $O(n)$  links: overall message complexity:  $O(n^2)$
-

# Dijkstra's algorithm: Oscillations


- When link costs depend on traffic volume, **route oscillations** possible
- Sample scenario:
  - Routing to destination a, traffic entering at d, c, e with rates 1, e ( $<1$ ), 1
  - Link costs are directional, and volume-dependent




initially



given these costs,  
find new routing....  
resulting in new  
costs



given these costs,  
find new routing....  
resulting in new costs



given these costs,  
find new routing....  
resulting in new costs

# Distance Vector Routing Protocol

---

# Distance vector algorithm

Based on **Bellman-Ford** (BF) equation (dynamic programming):

Bellman-Ford equation

Let  $D_x(y)$ : cost of least-cost path from  $x$  to  $y$ .

Then:


$$D_x(y) = \min_v \{ c_{x,v} + D_v(y) \}$$

$v$ 's estimated least-cost-path cost to  $y$

min taken over all neighbors  $v$  of  $x$       direct cost of link from  $x$  to  $v$

# Bellman-Ford Example

Suppose that  $u$ 's neighboring nodes,  $x, v, w$ , know that for destination  $z$ :



Bellman-Ford equation says:

$$\begin{aligned} D_u(z) &= \min \{ c_{u,v} + D_v(z), \\ &\quad c_{u,x} + D_x(z), \\ &\quad c_{u,w} + D_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad 1 + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$

node achieving minimum ( $x$ )  
is next hop on estimated  
least-cost path to  
destination ( $z$ )

# Distance vector algorithm

key idea:

- from time-to-time, each node sends its own distance vector estimate to neighbors
- when  $x$  receives new DV estimate from any neighbor, it updates its own DV using B-F equation:

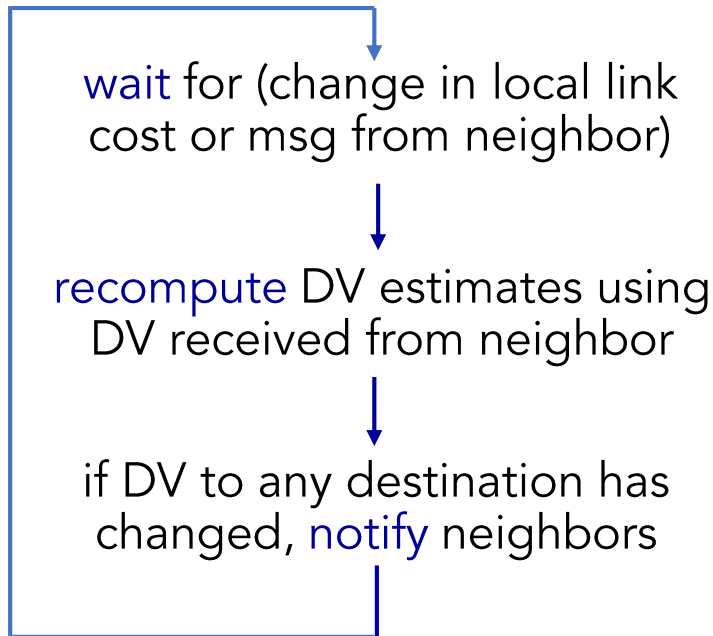
$$D_x(y) \leftarrow \min_v \{c_{x,v} + D_v(y)\} \text{ for each node } y \in N$$

- under minor, natural conditions, the estimate  $D_x(y)$  converge to the actual least cost  $d_x(y)$
-



# Distance vector algorithm:

each node:



**iterative, asynchronous:** each local iteration caused by:

- local link cost change
- DV update message from neighbor

**distributed, self-stopping:** each node notifies neighbors only when its DV changes

- neighbors then notify their neighbors – only if necessary
- no notification received, no actions taken!

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x  
table

		cost to		
		x	y	z
from	x	0	2	7
	y	$\infty$	$\infty$	$\infty$
	z	$\infty$	$\infty$	$\infty$


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

node y  
table

		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	2	0	1
	z	$\infty$	$\infty$	$\infty$

node z  
table

		cost to		
		x	y	z
from	x	$\infty$	$\infty$	$\infty$
	y	$\infty$	$\infty$	$\infty$
	z	7	1	0



time

$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x  
table

		cost to		
		x	y	z
from	x	0	2	7
	y	∞	∞	∞
	z	∞	∞	∞

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

node y  
table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	2	0	1
	z	∞	∞	∞

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	7	1	0


		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0

node z  
table

		cost to		
		x	y	z
from	x	∞	∞	∞
	y	∞	∞	∞
	z	7	1	0

		cost to		
		x	y	z
from	x	0	2	7
	y	2	0	1
	z	3	1	0

		cost to		
		x	y	z
from	x	0	2	3
	y	2	0	1
	z	3	1	0




time

# Distance vector: link cost changes

## link cost changes:

- ❖ node detects local link cost change
- ❖ updates routing info, recalculates distance vector
- ❖ if DV changes, notify neighbors



“good  
news  
travels  
fast”

$t_0$ : y detects link-cost change, updates its DV, informs its neighbors.


$t_1$ : z receives update from y, updates its table, computes new least cost to x, sends its neighbors its DV.

$t_2$ : y receives z's update, updates its distance table. y's least costs do not change, so y does not send a message to z.

# Distance vector: link cost changes

## link cost changes:

- ❖ node detects local link cost change
- ❖ **bad news travels slow** - “count to infinity” problem!
- ❖ 44 iterations before algorithm stabilizes



# Distance vector: link cost changes

## link cost changes:

- ❖ node detects local link cost change
- ❖ **bad news travels slow** - “count to infinity” problem!

A	B	C	D	E	
•	•	•	•	•	
	1	2	3	4	Initially
	3	2	3	4	After 1 exchange
	3	4	3	4	After 2 exchanges
	5	4	5	4	After 3 exchanges
	5	6	5	6	After 4 exchanges
	7	6	7	6	After 5 exchanges
	7	8	7	8	After 6 exchanges
		⋮			
	•	•	•	•	

## poisoned reverse:

- ❖ If Z routes through Y to get to X :
  - Z tells Y its (Z's) distance to X is infinite (so Y won't route to X via Z)
- ❖ will this completely solve count to infinity problem?

# Comparison of LS and DV algorithms

## message complexity

LS:  $n$  routers,  $O(n^2)$  messages sent

DV: exchange between neighbors;  
convergence time varies

## speed of convergence

LS:  $O(n^2)$  algorithm,  $O(n^2)$  messages

- may have oscillations

DV: convergence time varies

- may have routing loops
- count-to-infinity problem

robustness: what happens if router malfunctions, or is compromised?

LS:

- router can advertise incorrect link cost
- each router computes only its own table

DV:

- DV router can advertise incorrect path cost ("I have a really low cost path to everywhere"): black-holing
  - each router's table used by others: error propagate thru network
-

# THANK YOU

QUESTIONS???

---