

# Computer Networks I

## Application Layer

Amitangshu Pal

Computer Science and Engineering

IIT Kanpur

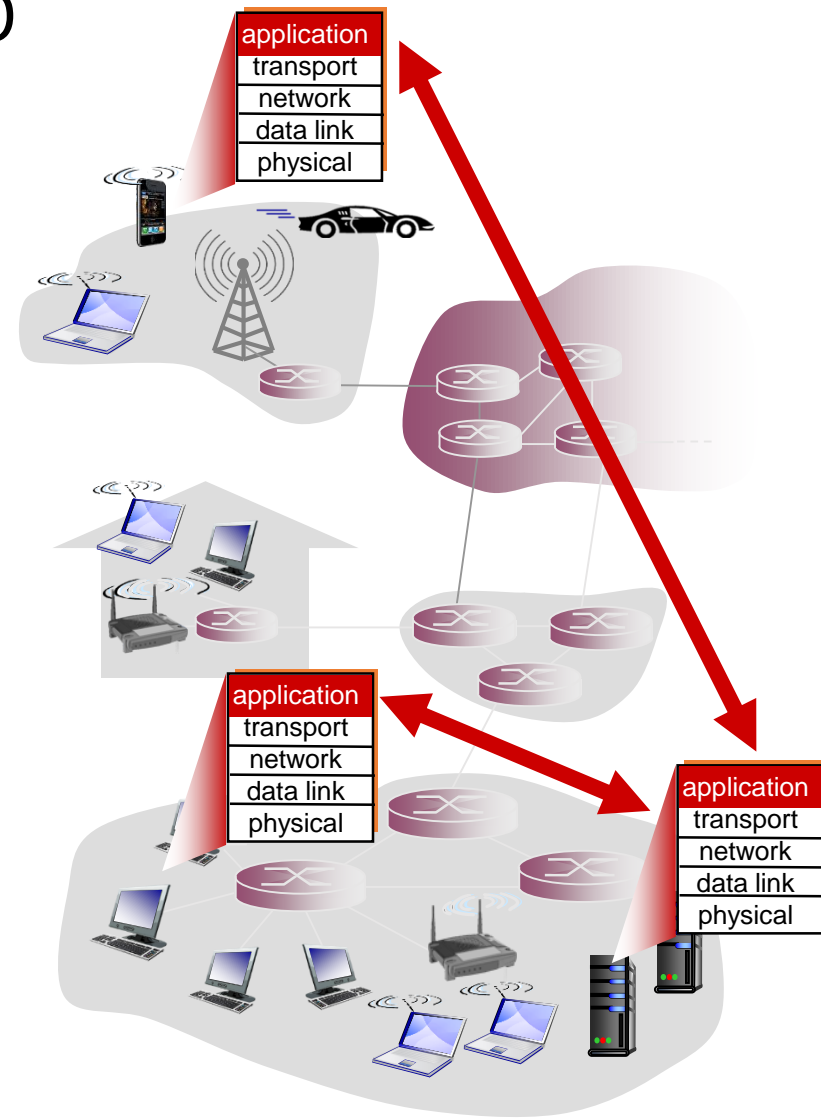
# Creating a network app

## write programs that:

- run on (different) *end systems*
- communicate over network
- e.g., web server software communicates with browser software

## no need to write software for network-core devices

- network-core devices do not run user applications
- applications on end systems allows for rapid app development, propagation




# Application architectures

possible structure of applications:

- client-server
  - peer-to-peer (P2P)
-

# Client-server architecture



## server:


- always-on host
- permanent IP address
- data centers for scaling

## clients:

- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not communicate directly with each other

# P2P architecture

- *no* always-on server
- arbitrary end systems directly communicate
- peers request service from other peers, provide service in return to other peers
  - *self scalability* – new peers bring new service capacity, as well as new service demands
- peers are intermittently connected and change IP addresses
  - complex management



# Internet apps: types and requirements

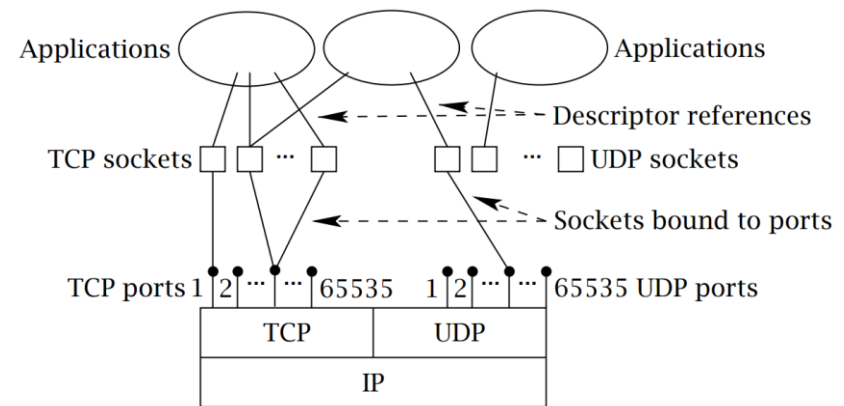
application	data loss	throughput	time sensitive
file transfer	no loss	elastic	no
e-mail	no loss	elastic	no
Web documents	no loss	elastic	no
real-time audio/video	loss-tolerant	audio: 5kbps-1Mbps video: 10kbps-5Mbps	yes, 100' s msec
stored audio/video	loss-tolerant	same as above	
interactive games	loss-tolerant	few kbps up	yes, few secs
text messaging	no loss	elastic	yes, 100' s msec yes and no

# Internet apps: application, transport protocols

<b>application</b>	<b>application layer protocol</b>	<b>underlying transport protocol</b>
e-mail	SMTP [RFC 2821]	TCP
remote terminal access	Telnet [RFC 854]	TCP
Web	HTTP [RFC 2616]	TCP
file transfer	FTP [RFC 959]	TCP
streaming multimedia	HTTP (e.g., YouTube), RTP [RFC 1889]	TCP or UDP
Internet telephony	SIP, RTP, proprietary (e.g., Skype)	TCP or UDP

# Internet apps: common port numbers

Port Number	Protocol	Application
20	TCP	FTP data
21	TCP	FTP control
22	TCP	SSH
23	TCP	Telnet
25	TCP	SMTP
53	UDP, TCP <sup>1</sup>	DNS
67	UDP	DHCP Server
68	UDP	DHCP Client
69	UDP	TFTP
80	TCP	HTTP (WWW)
110	TCP	POP3
161	UDP	SNMP





# Web and HTTP

---

# Web and HTTP

First, a review...

- **web page** consists of **objects**
- object can be HTML file, JPEG image, Java applet, audio file,...
- web page consists of **base HTML-file** which includes **several referenced objects**
- each object is addressable by a **URL**, e.g.,

`www.someschool.edu/someDept/pic.gif`

---

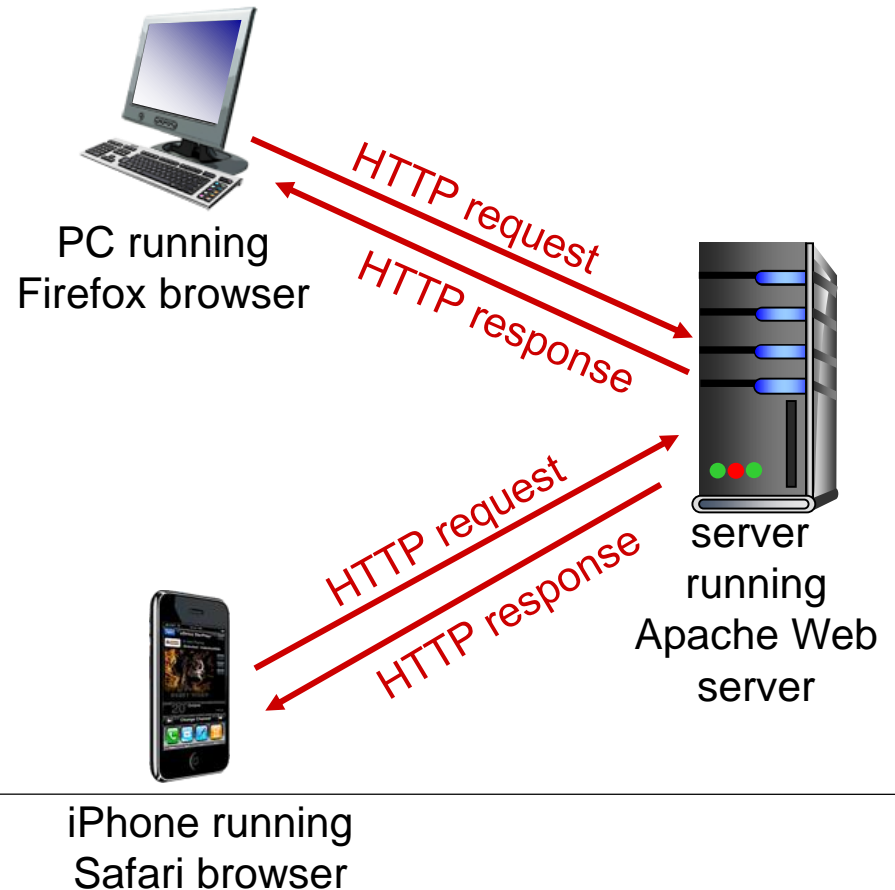
host name

path name

# HTTP overview

## HTTP: hypertext transfer protocol

- Web's application layer protocol
- client/server model
  - **client**: browser that requests, receives, (using HTTP protocol) and “displays” Web objects
  - **server**: Web server sends (using HTTP protocol) objects in response to requests



# HTTP overview

## uses TCP:

- client initiates TCP connection (creates socket) to server, port 80
  - server accepts TCP connection from client
  - HTTP messages (application-layer protocol messages) exchanged between browser (HTTP client) and Web server (HTTP server)
- 
- TCP connection closed

## HTTP is “stateless”

- server maintains no information about past client requests

# HTTP connections

## non-persistent HTTP

- at most one object sent over TCP connection
    - connection then closed
  - downloading multiple objects required multiple connections
- 

## persistent HTTP

- multiple objects can be sent over single TCP connection between client, server

# Non-persistent HTTP

suppose user enters URL:

`www.someSchool.edu/someDepartment/home.index`

(contains text,  
references to 10  
jpeg images)

1a. HTTP client initiates TCP connection to HTTP server (process) at `www.someSchool.edu` on port 80

1b. HTTP server at host `www.someSchool.edu` waiting for TCP connection at port 80. “accepts” connection, notifying client

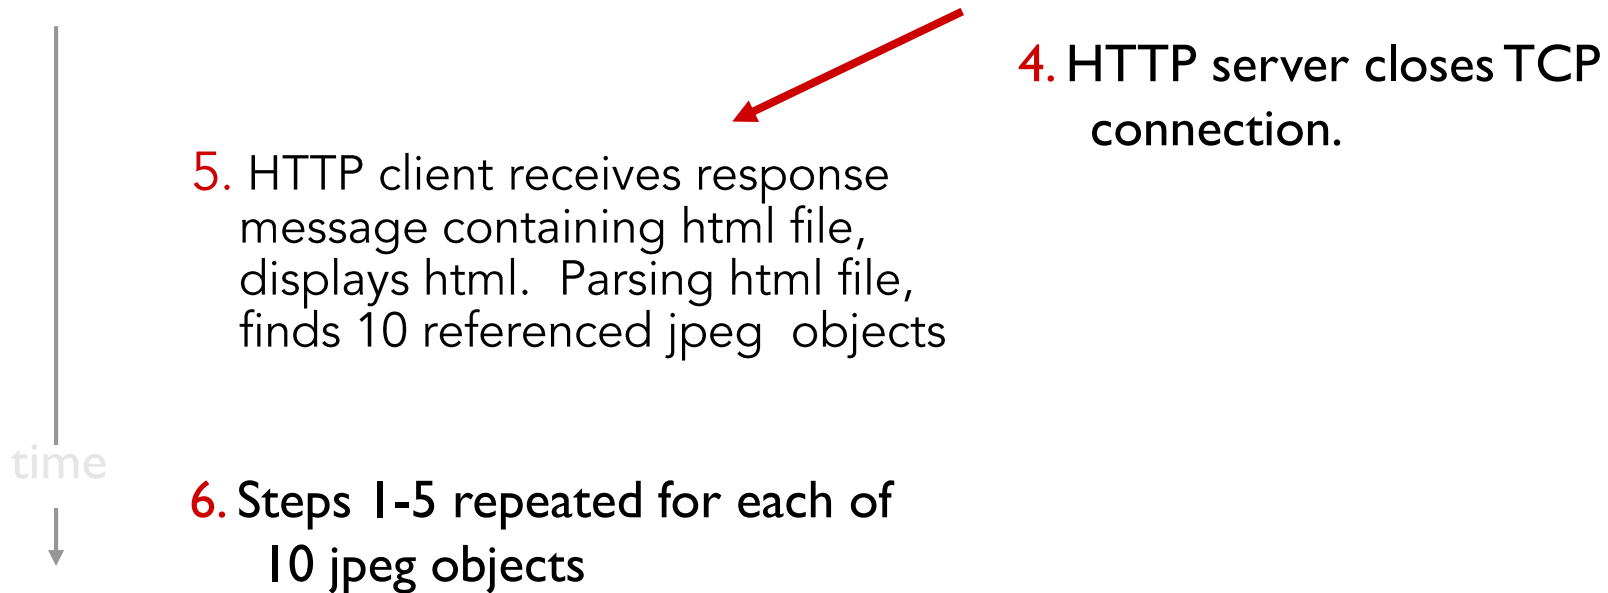
2. HTTP client sends HTTP *request message* (containing URL) into TCP connection socket. Message indicates that client wants object `someDepartment/home.index`

3. HTTP server receives request message, forms *response message* containing requested object, and sends message into its socket

time




# Non-persistent HTTP (cont.)



# Non-persistent HTTP: response time

## HTTP response time:

- one RTT to initiate TCP connection
- one RTT for HTTP request and first few bytes of HTTP response to return
- file transmission time
- non-persistent HTTP response time =  
 $2\text{RTT} + \text{file transmission time}$





# Persistent HTTP

## *non-persistent HTTP issues:*

- requires 2 RTTs per object
- OS overhead for *each* TCP connection
- browsers often open parallel TCP connections to fetch referenced objects

## *persistent HTTP:*

- server leaves connection open after sending response
  - subsequent HTTP messages between same client/server sent over open connection
  - client sends requests as soon as it encounters a referenced object
-

# User-server state: cookies

many Web sites use cookies

four components:

- 1) cookie header line of HTTP response message
  - 2) cookie header line in next HTTP request message
  - 3) cookie file kept on user's host, managed by user's browser
  - 4) back-end database at Web site
- 

example:

- An user visits specific e-commerce site for first time
- when initial HTTP requests arrives at site, site creates:
  - unique ID
  - entry in backend database for ID

# User-server state: cookies


client



server



cookie file



usual http request msg

usual http response  
**set-cookie: 1678**

usual http request msg  
**cookie: 1678**

usual http response msg


usual http request msg  
**cookie: 1678**

usual http response msg

Amazon server  
creates ID  
1678 for user

create  
entry

backend  
database




cookie-  
specific  
action

access

access

cookie-  
specific  
action

one week later:



# Cookies (continued)

what cookies can be used for:

- authorization
  - shopping carts
  - recommendations
  - user session state (Web e-mail)
-

# Web Caches

---


# Caching example:

## *assumptions:*

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: 1.54 Mbps

## *consequences:*

- LAN utilization: 15%
- access link utilization = **99%**
- total delay = Internet delay + access delay + LAN delay  
= 2 sec + minutes + usecs




# Caching example: fatter access link

*assumptions:*

- avg object size: 100K bits
- avg request rate from browsers to origin servers: 15/sec
- avg data rate to browsers: 1.50 Mbps
- RTT from institutional router to any origin server: 2 sec
- access link rate: ~~1.54 Mbps~~ → 154 Mbps

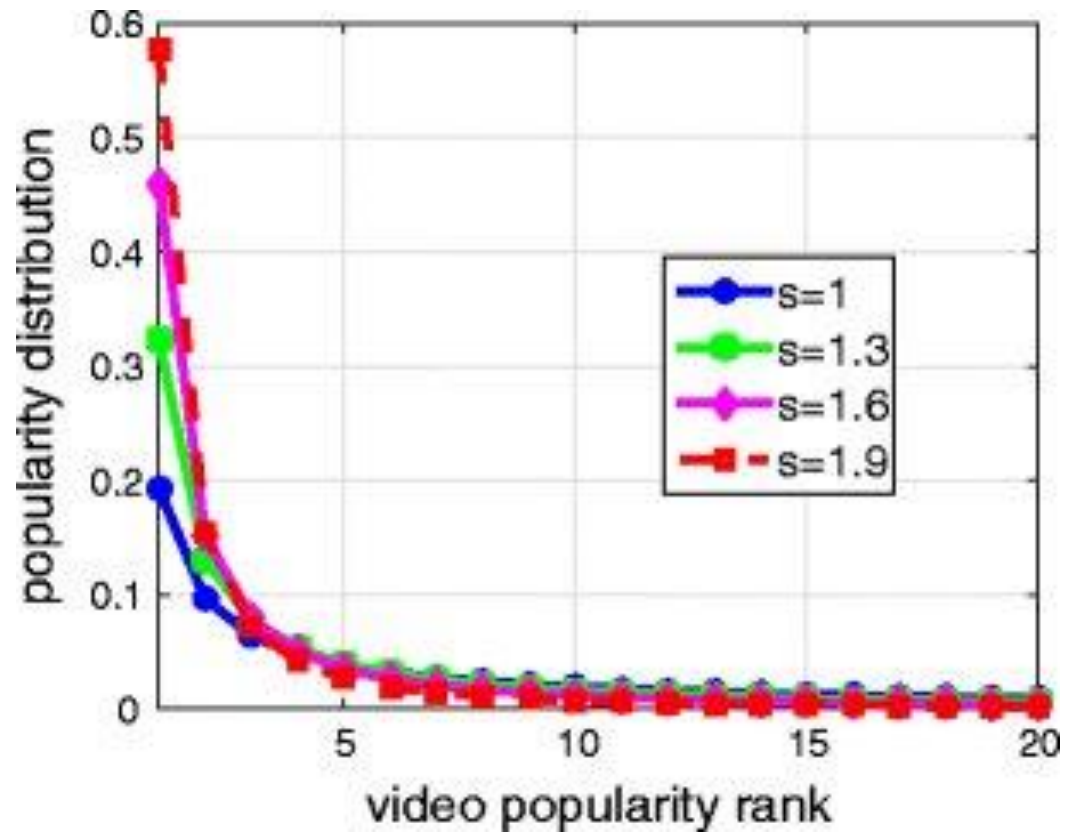
*consequences:*

- LAN utilization: 15%
- access link utilization = ~~99%~~ → 9.9%
- total delay = Internet delay + access  
delay + LAN delay  
= 2 sec + ~~minutes~~ → usecs  
msecs



**Cost:** increased access link speed (not cheap!)

# Internet content popularity




Src: [https://www.researchgate.net/publication/299518874\\_Fast-Start\\_Video\\_Delivery\\_in\\_Future\\_Internet\\_Architectures\\_with\\_Intra-domain\\_Caching/figures?lo=1](https://www.researchgate.net/publication/299518874_Fast-Start_Video_Delivery_in_Future_Internet_Architectures_with_Intra-domain_Caching/figures?lo=1)



# Caching example: install local cache

## *Calculating access link utilization, delay with cache:*

- suppose cache hit rate is 0.4
  - 40% requests satisfied at cache, 60% requests satisfied at origin
- access link utilization:
  - 60% of requests use access link
- data rate to browsers over access link  
 $= 0.6 * 1.50 \text{ Mbps} = .9 \text{ Mbps}$ 
  - utilization  $= 0.9 / 1.54 = .58$
- total delay
  - $= 0.6 * (\text{delay from origin servers}) + 0.4 * (\text{delay when satisfied at cache})$
  - $= 0.6 (2.01) + 0.4 (\sim \text{msecs}) = \sim 1.2 \text{ secs}$
  - less than with 154 Mbps link (and cheaper too!)

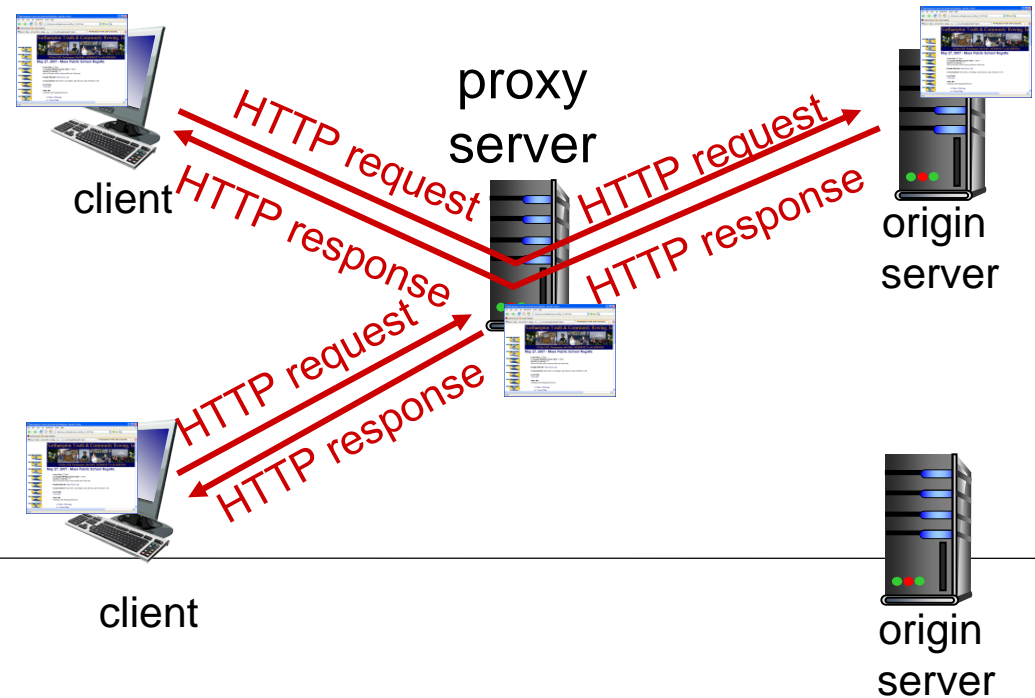


# Web caches (proxy server)

**goal:** satisfy client request without involving origin server

## why Web caching?

- reduce response time for client request
- reduce traffic on an institution's access link
- Internet dense with caches: enables “poor” content providers to effectively deliver content



# Conditional GET

- **Goal:** don't send object if cache has up-to-date cached version

- no object transmission delay
- lower link utilization

- **cache:** specify date of cached copy in HTTP request

**If-modified-since:**  
**<date>**

- **server:** response contains no object if cached copy is up-to-date:

**HTTP/1.0 304 Not Modified**

client



server



HTTP request msg  
**If-modified-since: <date>**

object  
not  
modified  
before  
<date>

HTTP response  
**HTTP/1.0  
304 Not Modified**

-----

HTTP request msg  
**If-modified-since: <date>**

object  
modified  
after  
<date>


HTTP response  
**HTTP/1.0 200 OK  
<data>**

# Electronic Mail

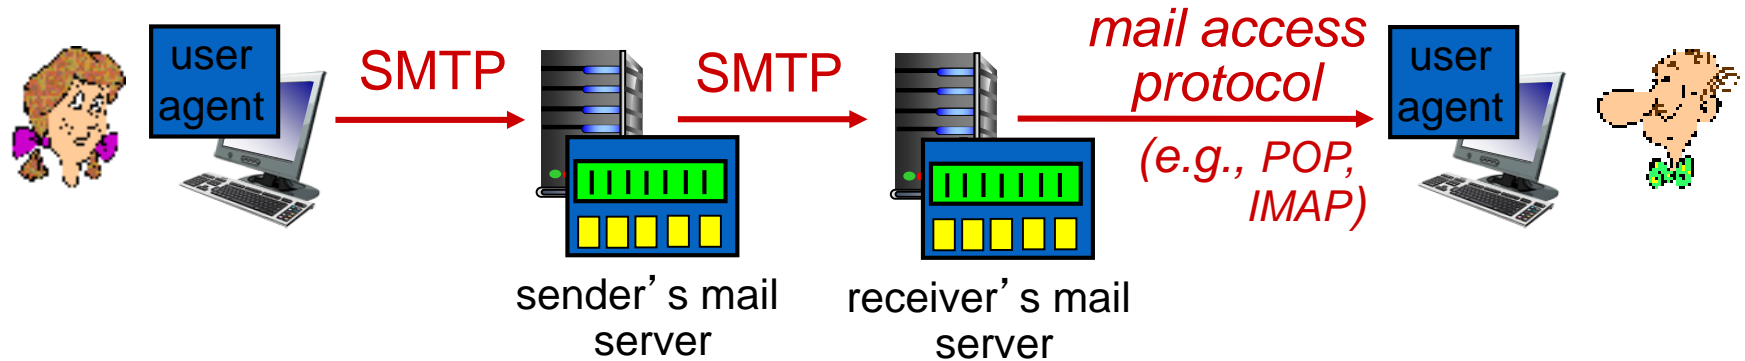
---

# Electronic mail

- **SMTP protocol** between mail servers to send email messages
  - client: sending mail server
  - “server”: receiving mail server



# Mail access protocols



- **SMTP**: delivery/storage to receiver's server
- mail access protocol: retrieval from server
  - **POP**: Post Office Protocol [RFC 1939]: authorization, download
  - **IMAP**: Internet Mail Access Protocol [RFC 1730]: more features, including manipulation of stored messages on server
  - **HTTP**: gmail, Hotmail, Yahoo! Mail, etc.

DNS

---

# DNS: domain name system

## Internet hosts, routers:

- IP address (32 bit) - used for addressing datagrams
- “name”, e.g., `www.yahoo.com` - used by humans

## DNS services


- hostname to IP address translation

## Domain Name System:

- distributed database implemented in hierarchy of many name servers
  - application-layer protocol: hosts, name servers communicate to resolve names (address/name translation)
-



# DNS: a distributed, hierarchical database

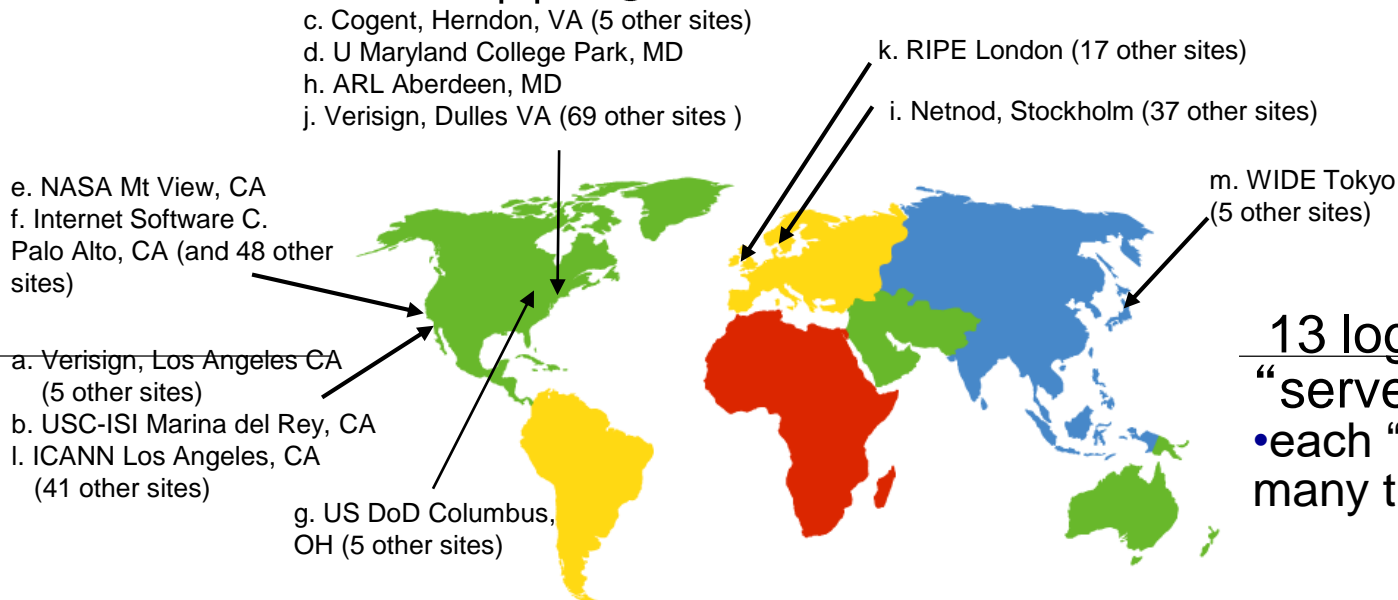


client wants IP for [www.amazon.com](http://www.amazon.com):

- client queries root server to find com DNS server
- client queries .com DNS server to get amazon.com DNS server
- client queries amazon.com DNS server to get IP address for [www.amazon.com](http://www.amazon.com)

# DNS: root name servers

- contacted by local name server that can not resolve name
- root name server:
  - contacts authoritative name server if name mapping not known
  - gets mapping
  - returns mapping to local name server



---

13 logical root name  
“servers” worldwide  
• each “server” replicated  
many times

# TLD, authoritative servers

## top-level domain (TLD) servers:

- responsible for com, org, net, edu, aero, jobs, museums, and all top-level country domains, e.g.: uk, fr, ca, jp
- Network Solutions maintains servers for .com TLD
- Educause for .edu TLD

## authoritative DNS servers:

- organization's own DNS server(s), providing authoritative hostname to IP mappings for organization's named hosts
  - can be maintained by organization or service provider
-

# Local DNS name server


- does not strictly belong to hierarchy
  - each ISP (residential ISP, company, university) has one
    - also called “default name server”
  - when host makes DNS query, query is sent to its local DNS server
    - has local cache of recent name-to-address translation pairs (but may be out of date!)
    - acts as proxy, forwards query into hierarchy
-

# DNS name resolution example

- host at cis.poly.edu wants IP address for gaia.cs.umass.edu

## *iterated query:*

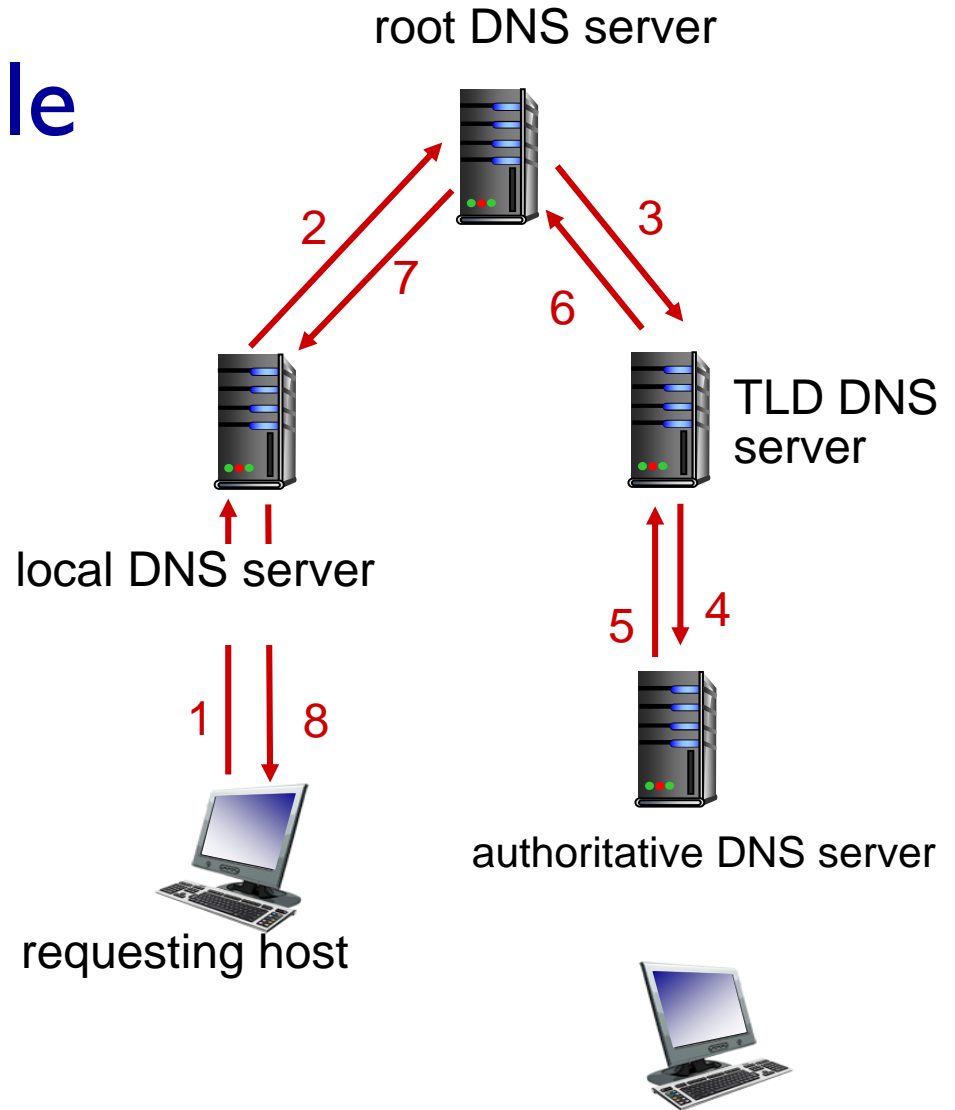
- contacted server replies with name of server to contact
- “I don’t know this name, but ask this server”



# DNS name resolution example

## *recursive query:*

- puts burden of name resolution on contacted name server
- heavy load at upper levels of hierarchy?



# DNS: caching, updating records

- once (any) name server learns mapping, it **caches** mapping
    - cache entries timeout (disappear) after some time (TTL)
    - TLD servers typically cached in local name servers
      - thus root name servers not often visited
-

# THANK YOU

QUESTIONS???

---